Spectral indices and incidental amplitudes in pixel-based Bayesian component separation: a tempered sampler, the reference frequency of a flat prior, and an honest error for the profile likelihood

T. L. Svalheim

Draft of November 11, 2026. Internal, not for distribution. Every number carries the identifier of the experiment record it comes from, in the form E035; the records are in `nextgen/experiments/`. Red text marks what is still pending.

Abstract

In pixel-based Bayesian component separation a few spectral indices share the data with a very large number of per-pixel amplitudes. We study what three common ways of removing the amplitudes do to the index. (1) A Metropolis move that proposes an index, redraws the amplitudes and accepts on the χ^2 of the new pair is, on its own, a pseudo-marginal kernel for the profile posterior; inside a Gibbs sweep that also redraws the amplitudes it becomes Monte Carlo within Metropolis and, for small steps, converges to the profile posterior raised to the power $1/T$, with $T = \frac{1}{2} + 1/(4c_k) \simeq \sqrt{\pi k/8} + \frac{1}{2}$ for k amplitudes. The index error from such a chain is too wide by $0.79\,k^{1/4}$. We derive the law and its form for several moves per amplitude draw and for inexact stored amplitudes, and confirm it against an exact-kernel calculation and simulated chains within their Monte Carlo errors of 0.4 percent. (2) Moving the frequency at which a flat amplitude prior is stated multiplies the index prior by an exponential with rate $n \ln(\nu'_{\rm ref}/\nu_{\rm ref})$, n the number of amplitudes. For one component the index is shifted by $\ln(\nu_{\rm ref}/\nu_{\rm piv})$ times the index variance carried by one amplitude, a shift that does not average down over pixels, changes sign across a pivot frequency set by the bands and by the prior on the CMB, and has opposite signs for references of 23 and 30 GHz with the WMAP and Planck LFI bands. (3) The profile likelihood has an exactly unbiased score, but its curvature is too sharp by $\sqrt{1 + \mathsf{T}/\mathsf{J}}$, where T/J is the inverse mean squared signal-to-noise ratio of the amplitudes; the natural plug-in sandwich correction is biased the wrong way, and a debiased form gives nominal coverage over four decades of T/J. For one component and a uniform noise pattern a zero-mean amplitude prior with a free scale is the same estimator; otherwise it is biased by up to 0.6 standard deviations. We then apply the estimators side by side to the WMAP nine-year K to V polarization maps with their dense noise covariances and a ΛCDM prior on the CMB. On the same maps, regions and noise model the high-latitude polarized synchrotron index is -3.08 with the profile likelihood, -3.09 with a flat amplitude prior stated at 23 GHz and -2.92 with the same prior stated at 30 GHz: 12.6 statistical standard deviations from a change of convention, as the identity predicts to 9 percent. With a free CMB amplitude the same change of convention moves it by 0.8, again as predicted. A white prior on the CMB with the right total variance leaves the index unbiased on simulated ΛCDM skies and makes its error bar too small by a quarter to a third. We do not quote an index. The scatter of the score over patches of sky shows that the high-latitude region and the Galactic plane are not one index each, and gives errors of 0.04 and 0.10 where the model gives 0.012 and 0.013; the index also moves by 0.02 to 0.11 with the angular scales that inform it, by 0.01 to 0.07 with the prior on the CMB and by up to 0.03 for 0.05 in the fixed dust index; and the official Planck 30 GHz polarization maps do not fit one power law with WMAP K band. We also report a silent memory fault in a vendor Cholesky routine that sat under every low-resolution fit of this size on our machine, how we found it, and that no number changed.

1 Introduction

Bayesian component separation in the pixel domain fits, in every pixel, the amplitude of each sky component, and for the whole sky or for large regions a small number of spectral parameters

that say how each component scales with frequency. In the Gibbs samplers that the Planck, BeyondPlanck and Cosmoglobe analyses are built on (Eriksen et al., 2008; Svalheim et al., 2023; Watts et al., 2023), the two kinds of parameter are updated in turn. The amplitudes are many and each is poorly determined; the spectral indices are few and are what the cleaning of large-scale polarization, and therefore the optical depth to reionization, depends on. The published polarized synchrotron indices from the same public maps span -2.9 to -3.3 depending on region, band set and method (Dunkley et al., 2009; Fuskeland et al., 2014; Svalheim et al., 2023; Weiland et al., 2022), which is larger than their quoted errors.

This paper is about the statistical step that sits between the maps and such an index: how the amplitudes are removed. That step has three places where it can go wrong without any error message, and we treat them in turn.

The sampler (Sect. 3). A move that proposes a new index, re-solves the amplitudes and accepts on the χ^2 of the new pair looks like a way to move both together. We show that on its own it is a valid pseudo-marginal kernel for the profile posterior, that the amplitude step of the Gibbs sweep turns it into Monte Carlo within Metropolis, and that the chain then converges to the profile posterior raised to the power $1/T$ with $T = \sqrt{\pi k/8} + \frac{1}{2}$ for k amplitudes. The index error bar from such a chain is too wide by $0.79\,k^{1/4}$, a factor 16 for the 1.9×10^5 amplitudes of an $N_{\rm side} = 64$ analysis, with an acceptance rate that stays near one half.

The prior (Sect. 4). With flat priors on the amplitudes, the standard choice, the index posterior contains the volume of amplitude space. This is known to bias indices (Eriksen et al., 2008; Dunkley et al., 2009; Alonso et al., 2017). We show that it is exactly equivalent to an exponential prior on the index whose rate is the number of amplitudes times the logarithm of the reference frequency at which the amplitudes are defined, give the resulting shift in closed form, and show that its sign flips between the two reference frequencies in common use, 23 and 30 GHz, for the WMAP and LFI bands. The shift does not shrink as pixels are added.

The error bar (Sect. 5). The profile likelihood has neither problem and its score is exactly unbiased, but its curvature understates the error where the amplitudes are faint. We give the correction in closed form, show that the obvious way to evaluate it on real data is biased in the wrong direction, give the version that is not, and show its coverage over four decades of the controlling ratio. A zero-mean amplitude prior with a free scale turns out to be the same estimator when the noise pattern is uniform, and a biased one when it is not.

Sect. 6 applies the estimators side by side to the WMAP nine-year K to V band polarization maps with their released dense noise covariances, so that the difference between their answers is the difference between the estimators. It then lists what else moves the index on these maps, by how much, and why we stop short of quoting one. Sect. 7 states where the recommended estimator stops being enough, and Sect. 8 describes the implementation and its tests. All results are for maps-only analyses at low resolution with amplitudes that enter linearly; we do not treat the time-ordered data.

2 Model, notation and the index targets in use

2.1 Data model

The data are maps d_b in bands $b = 1, \ldots, n_{\rm b}$, each a vector over elements p (pixels, and Stokes parameters in polarization), in thermodynamic units. The parametric model is

$$d_b = \sum_c m_{b,c}(\beta)\, a_c + n_b, \qquad n_b \sim \mathcal{N}(0, \mathsf{N}_b), \tag{1}$$

with one amplitude $a_c(p)$ per component c and element, a mixing coefficient $m_{b,c}$ that is the component spectrum averaged over the bandpass of band b, and spectral parameters β that are shared by many elements: one value for the sky, one per region, or a field at lower resolution than

the amplitudes. Stacking bands and components, $d = \mathsf{M}(\beta)\, a + n$ with a of length k. We write

$$\mathsf{A} = \mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}\mathsf{M}, \qquad \hat{a} = \mathsf{A}^{-1}\mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}d, \qquad \mathsf{P} = \mathsf{N}^{-1} - \mathsf{N}^{-1}\mathsf{M}\mathsf{A}^{-1}\mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}, \tag{2}$$

all functions of β, so that

$$\chi^2(\beta, a) = \chi^2_{\min}(\beta) + (a - \hat{a})^{\mathsf{T}}\mathsf{A}\,(a - \hat{a}), \qquad \chi^2_{\min}(\beta) = d^{\mathsf{T}}\mathsf{P}d. \tag{3}$$

For the power-law component that carries the index, $m_b(\beta)$ is the band average of $(\nu/\nu_{\mathrm{ref}})^\beta$ in brightness temperature, converted to thermodynamic units, so that the amplitude is the brightness temperature at the reference frequency ν_{ref}. The reference frequency is a convention: it changes the meaning of a and not the model.

The amplitudes are what statisticians call incidental parameters (Neyman & Scott, 1948): their number grows with the data, and the parameters of interest, the indices, are few. Everything in this paper is a consequence of how the incidental parameters are removed.

2.2 Four ways to remove the amplitudes

Joint posterior with flat amplitude priors. With $p(a) \propto 1$ the joint posterior is $p(\beta)\exp[-\chi^2(\beta, a)/2]$ and integrating Eq. (3) over a gives the index marginal

$$\mathcal{L}_{\mathrm{flat}}(\beta) \propto \exp\!\left[-\tfrac{1}{2}\chi^2_{\min}(\beta)\right] \det \mathsf{A}(\beta)^{-1/2}. \tag{4}$$

A Gibbs sampler that draws $a \mid \beta$ with its fluctuation and then moves β on χ^2 at fixed a has exactly this marginal for β, whether or not it mixes well enough to show it (checked numerically and re-derived independently, E012). The determinant is the volume of amplitude space allowed at a given β. It is an implicit prior on the index and the subject of Sect. 4.

Profile likelihood. Maximizing over the amplitudes instead of integrating gives

$$\mathcal{L}_{\mathrm{prof}}(\beta) \propto \exp\!\left[-\tfrac{1}{2}\chi^2_{\min}(\beta)\right], \tag{5}$$

the spectral likelihood of Stompor et al. (2009) (their Eq. 9; their Eq. 10 is the marginal form). As a Bayesian target it is the marginal under the amplitude prior $p(a \mid \beta) \propto \det \mathsf{A}(\beta)^{1/2}$, a Jeffreys prior for the amplitude block conditional on the index. It is not the marginal under the joint Jeffreys prior of (a, β), which depends on a.

Gaussian amplitude prior. With $a \sim \mathcal{N}(\mu, \mathsf{S})$ the marginal keeps the form of Eq. (4) with $\mathsf{A} \to \mathsf{A} + \mathsf{S}^{-1}$, the prior mean entering $\chi^2_{\min}$, and a factor $\det \mathsf{S}^{-1/2}$ that matters when S is itself inferred. We use two cases. A zero-mean prior with one free scale per index region, $a_{\mathrm{s}}(p) \sim \mathcal{N}(0, \tau_r^2)$ for p in region r, with τ_r maximized jointly with the index (a prior flat in $\ln \tau_r$ needs a lower cutoff to be proper, since the likelihood tends to a constant as $\tau_r \to 0$). And, for intensity, a prior whose mean and scale are tied to a template, with hyperparameters inferred. Under a change of reference frequency τ_r rescales by $(\nu'_{\mathrm{ref}}/\nu_{\mathrm{ref}})^{\beta_r}$, and maximization is invariant under that rescaling, so the first case does not depend on the reference frequency.

Re-solving the amplitudes inside the index move. A fourth route is procedural: propose β', solve for new amplitudes at β', and accept on the χ^2 of the new pair. What this samples depends on how the new amplitudes are produced, and is the subject of Sect. 3.

2.3 Information and score

Let $\mathsf{M}_r = \partial\mathsf{M}/\partial\beta_r$ for index parameter r. For data that follow Eq. (1) with true amplitudes a, the profile score $s_r = -\frac{1}{2}\,\partial\chi^2_{\min}/\partial\beta_r$ has zero mean and covariance

$$\mathrm{Cov}(s_r, s_{r'}) = \mathsf{J}_{rr'} + \mathsf{T}_{rr'}, \qquad \mathsf{J}_{rr'} = (\mathsf{M}_r a)^{\mathsf{T}}\mathsf{P}\,(\mathsf{M}_{r'} a), \qquad \mathsf{T}_{rr'} = \mathrm{tr}\Big(\mathsf{A}^{-1}\mathsf{M}_r^{\mathsf{T}}\mathsf{P}\,\mathsf{M}_{r'}\Big)\,, \tag{6}$$

while the expected curvature of $\frac{1}{2}\chi^2_{\min}$ is J alone (derivation in App. A). These statements are for flat priors on every amplitude; with a Gaussian prior on a component they hold on average over realizations of that component drawn from the prior, with its precision included in A. J is the Fisher information on the indices with the amplitudes profiled out. T does not depend on the signal: it is the part of the score variance that the noise in the fitted amplitudes contributes. For one power-law component, diagonal noise with weights $w_{pb} = 1/\sigma^2_{pb}$ and one region, these reduce to

$$\mathsf{J} = \sum_p u_p V_p, \qquad \mathsf{T} = \sum_p V_p, \qquad u_p = a_p^2 \sum_b w_{pb} m_b^2, \qquad V_p = \mathrm{Var}_{wm^2}(\ln\nu)_p, \tag{7}$$

with u_p the squared signal-to-noise ratio of amplitude p and V_p the variance of $\ln\nu$ over the bands under the weights $w_{pb}m_b^2$: the lever arm.

3 What the re-solve move converges to

3.1 The move

The move we analyse is natural, and a version of it was the spectral-index step of the per-pixel path in the development version of Commander4 that we started from (option `resolve_chisq` in our branch). Given the current (β, a):

1. propose β' from a symmetric random walk;
2. draw new amplitudes from their conditional at the proposed index, $a' \sim \mathcal{N}\big(\hat{a}(\beta'), \mathsf{A}(\beta')^{-1}\big)$;
3. accept (β', a') with probability $\min\{1, \exp[-(\chi^2(\beta', a') - \chi^2(\beta, a))/2]\, p(\beta')/p(\beta)\}$, with p the index prior.

The Gibbs sweep then continues with the amplitude group, which redraws $a \sim \mathcal{N}(\hat{a}(\beta), \mathsf{A}(\beta)^{-1})$, and possibly repeats the index move n times before that. The intent is to move the index and the amplitudes together and so avoid the slow mixing of the conditional index step. Write $X = (a - \hat{a})^{\mathsf{T}}\mathsf{A}(a - \hat{a})$, so that $\chi^2 = \chi^2_{\min}(\beta) + X$ by Eq. (3). For an amplitude draw from the conditional, $X \sim \chi^2_k$ whatever β is.

3.2 Alone, the move is exact for the profile posterior

Because $\mathrm{E}[\exp(-X/2)] = 2^{-k/2}$ does not depend on β, the quantity $\exp[-\chi^2(\beta, a)/2]$ at a fresh conditional draw is an unbiased estimator of $\mathcal{L}_{\mathrm{prof}}(\beta)$ up to a constant. Steps 1 to 3, with the accepted amplitudes kept as part of the state, are therefore a pseudo-marginal Metropolis-Hastings kernel (Beaumont, 2003; Andrieu & Roberts, 2009). Its extended target is

$$\pi(\beta, a) \propto p(\beta)\, \exp\big[-\tfrac{1}{2}\chi^2_{\min}(\beta)\big]\, \det \mathsf{A}(\beta)^{1/2} \exp\Big[-(a - \hat{a})^{\mathsf{T}}\mathsf{A}\,(a - \hat{a})\Big]\,, \tag{8}$$

whose index marginal is the profile posterior and whose amplitude conditional is $\mathcal{N}(\hat{a}, (2\mathsf{A})^{-1})$: the amplitudes carry half their posterior covariance. So the move does have a target, the target is the profile posterior, and the determinant never enters. This is why chains run with this move have means close to the profile posterior. Under Eq. (8), $X \sim \mathrm{Gamma}(k/2, 1)$ with mean $k/2$.

3.3 The amplitude step of the Gibbs sweep breaks it

The amplitude group redraws a with the full conditional covariance A^{-1}, which resets X to a χ^2_k variate with mean k. That is the conditional of the flat-prior joint posterior and not of Eq. (8). The two kernels share no invariant distribution. With one index move per amplitude draw, both the current and the proposed estimate are fresh at every index move, and the composite is Monte Carlo within Metropolis in the terminology of Andrieu & Roberts (2009), following O'Neill et al. (2000) and Beaumont (2003). Andrieu & Roberts (2009) note that the target is typically not the invariant density of such a chain and that the existence of an invariant distribution is not obvious in general; later work bounds the distance between the noisy and the exact chain without giving the stationary law (Alquier et al., 2016; Medina-Aguayo et al., 2016). Here it can be written down in the limit of small proposal steps.

3.4 The stationary law: a tempered profile posterior

With one move per amplitude draw the index chain is Markov on its own, with acceptance probability

$$\alpha(D) = \mathrm{E}\big[\min\{1, e^{-(D+Z)/2}\}\big], \qquad D = \chi^2_{\min}(\beta') - \chi^2_{\min}(\beta) - 2\ln\frac{p(\beta')}{p(\beta)}, \qquad Z = X' - X, \quad (9)$$

where X and X' are independent χ^2_k variates, so Z is symmetric with variance $4k$. In the limit of small proposal steps the chain is a diffusion with drift proportional to $\alpha'(0)$ and diffusion coefficient proportional to $\alpha(0)$, and its stationary density is

$$\pi(\beta) \propto \big\{p(\beta)\exp\big[-\tfrac{1}{2}\chi^2_{\min}(\beta)\big]\big\}^{1/T}, \qquad T = \frac{1}{2} + \frac{1}{4c_k}, \qquad c_k = \mathrm{E}\big[e^{-Z/2};\, Z > 0\big] = 2^{-(k+1)}\binom{k}{k/2}. \quad (10)$$

The formula is exact for every k and every shape of $\chi^2_{\min}$ under four conditions: the amplitude draws are exact conditional draws; the k amplitudes redrawn by the Gibbs step are the ones that enter the acceptance; the proposal is symmetric with a fixed covariance; and the proposal steps are small against the stationary width. $T(0) = 1$ recovers Metropolis. For large k,

$$T = \sqrt{\frac{\pi k}{8}} + \frac{1}{2} + O(k^{-1/2}), \quad (11)$$

so the index posterior is the profile posterior with its width multiplied by $\sqrt{T} \simeq 0.79\, k^{1/4}$. The stationary acceptance rate is $\frac{1}{2} + c_k$, close to one half and independent of the step size, which is why such a chain looks healthy.

Equation (10) does not need χ^2_k noise. For any symmetric acceptance noise Z whose density f_Z is flat over a few units around zero, $T \simeq \frac{1}{2} + 1/(8 f_Z(0))$, and for Gaussian noise

$$T \simeq \frac{1}{2} + \sqrt{\frac{\pi}{8}}\,\sigma_E, \qquad \sigma_E = \tfrac{1}{2}\,\mathrm{sd}(Z), \quad (12)$$

with σ_E the standard deviation of the noise in the estimated energy difference, equal to $1/\sqrt{2}$ times the standard deviation of the χ^2 at a fresh amplitude draw. The leading term is the effective temperature that Ball et al. (2003) derived for the sign rule (accept if and only if the noisy energy difference is negative; their Eq. 7), and Bardenet et al. (2017) obtain, without assuming Gaussian noise, a temperature proportional to the noise with a slightly different constant (their Sect. 6.3). That a Metropolis chain on a noisy energy difference is over-dispersed is noted by Nicholls et al. (2012), and Ceperley & Dewing (1999) give the exact correction for Gaussian noise of known variance. The Metropolis rule on the same noisy difference adds the $\frac{1}{2}$. Equation (12) is the form to use when the stored amplitudes are not the drawn ones, as below.

Table 1: Variance of the stationary index distribution of the composite chain in units of the profile posterior variance, one index move per amplitude draw, for $\chi^2_{\min} = \beta^2$ and exact χ^2_k refreshes. Columns: Eq. (10); its leading term; the stationary vector of the exact kernel on a grid of 1201 points, with the acceptance function computed by convolution of the χ^2_k density; and a simulation of 10^4 chains with a proposal step of $0.1\sqrt{T}$ (E035, independent review). The grid and the simulation sit 0.1 to 0.2 percent below the law because of the finite step. Controls: with no refresh, or with a refresh from Gamma$(k/2, 1)$, the variance is 1.000 ± 0.001 at $k = 4$, 8 and 16.

k	T, Eq. (10)	$\sqrt{\pi k/8}$	exact kernel	simulation	acceptance
10	2.532	1.982	2.524	2.525 ± 0.008	0.621
100	6.782	6.267	6.769	6.751 ± 0.024	0.538
1 000	20.32	19.82	20.29	20.26 ± 0.07	0.511
10 000	63.17	62.67	63.07	63.30 ± 0.23	0.503
100 000	198.67	198.17	198.35	199.39 ± 0.72	0.500
200 000	280.75	280.25	280.31	280.22 ± 1.01	0.500

Several index moves per amplitude draw. After the first accepted move the retained X is no longer fresh: it is biased low, and the next proposal has to beat it. In the large-k limit the accepted moves are the lower records of the sequence of fresh X values, the j-th proposal after a redraw is accepted with probability $1/(j+1)$, and the stationary law remains a power of the profile posterior with temperature $T_n = R_n T_1$,

$$R_n = \frac{H_{n+1} - 1}{\sqrt{\pi}\, N_\varphi(n)}, \qquad N_\varphi(n) = \sum_{j=1}^{n} j \int \varphi(t)^2 \left[1 - \Phi(t)\right]^{j-1} \mathrm{d}t, \tag{13}$$

where H_n is the harmonic number, φ and Φ are the standard normal density and distribution function, and $N_\varphi(n)$ is the sum over proposals of the expected density at the running minimum. $R_2 = 5/6$ exactly, and numerically $R_n = 0.833$, 0.647, 0.543, 0.467 and 0.357 for $n = 2$, 5, 10, 20 and 100: more index moves narrow the law only slowly, asymptotically as $(\ln n)^{-1/4}$ in width, and the $k^{1/4}$ scaling is untouched.

The repair that is exact and useless. Redrawing the amplitudes with half their covariance, $a \sim \mathcal{N}(\hat{a}, (2\mathsf{A})^{-1})$, is the conditional of Eq. (8) and makes the composite exact for the profile posterior. Its acceptance rate at fixed index is $2I_{1/3}(k/2, k/2)$, with I the regularized incomplete beta function, which decays as $e^{-0.059k}$: 6×10^{-4} at $k = 100$. The useful repairs evaluate Eq. (5) or Eq. (4) directly in the acceptance, which costs one extra log-determinant per proposal at most (Sect. 8).

3.5 Tests

Table 1 and Fig. 1 compare Eq. (10) with the exact kernel and with simulated chains. The law was derived twice, the second time by an independent reviewer who was given the algorithm and not our derivation. The two derivations agree, the reviewer's exact-kernel calculation reproduces every row of our own toy within 1.4 standard errors, and the controls return the profile posterior exactly, so the excess is caused by the full-covariance redraw and by nothing else. For proposal steps that are not small the kernel is no longer reversible with respect to a power of the profile posterior, but for $k \geq 10^4$ the variance stays within 8 percent of T at any step size (0.918 T for steps of four stationary widths or more; 0.907 T at $k = 100$).

The same toy reproduces, within 5 percent, four earlier measurements of the width of chains run with this move on simulated maps, with k from 450 to 3072 and one or four moves per amplitude draw (measured 3.9, 5.7, 3.2 and 4.4 profile widths; toy 4.0 to 4.1, 5.7, 3.0 and 4.3; E002, E005).

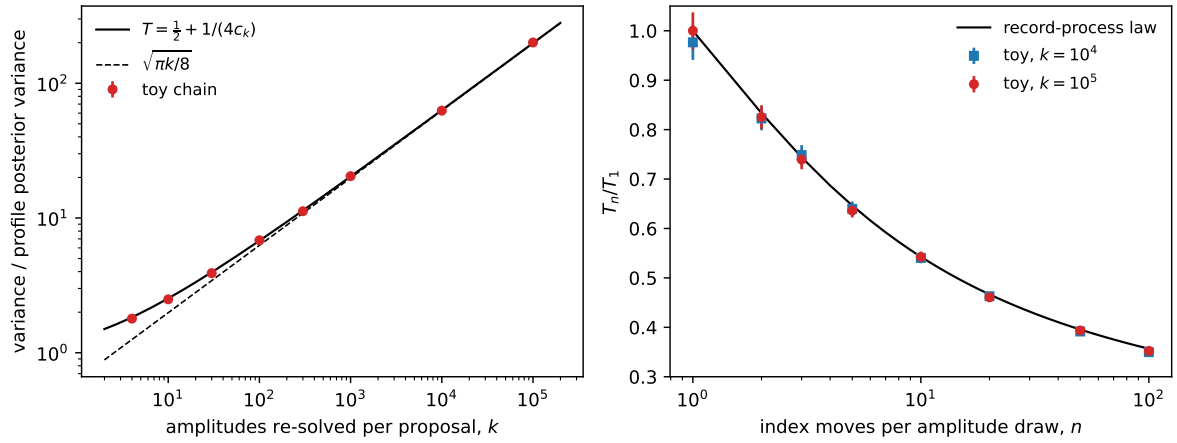


Figure 1: The stationary law of the re-solve move inside a Gibbs sweep. Left: temperature of the tempered profile posterior against the number of amplitudes re-solved per proposal, for one index move per amplitude draw: Eq. (10) (line), its leading term (dashed) and the toy chain (points). Right: the ratio T_n/T_1 for n index moves per amplitude draw: the record-process law of Eq. (13) (line) and the toy at $k = 10^4$ and 10^5 (points).

Real maps. On nside-64 intensity maps (17 bands from Haslam to Planck 857 GHz, 38 165 unmasked pixels, five components, so $k = 190\,825$) a chain with 20 index moves per amplitude draw, proposing the synchrotron and dust indices jointly with steps of 5.6 and 93 profile widths, was run twice. The widths were 17.0 and 16.4 profile widths for the synchrotron index (about 8 percent error each) and 16.4 and 16.6 for the dust index (2 percent each) (`E014`). For small steps Eq. (13) gives $\sqrt{R_{20}T_1} = 11.3$; a toy chain with two indices and the recorded step sizes gives 13.1 and 13.6. The measured widths are 20 to 30 percent above that. We can account for the excess after the fact, and say so plainly: the account uses one number measured from the log of the same chain. This code path stores the per-pixel amplitudes as spherical-harmonic coefficients to $\ell_{\max} = 3N_{\rm side} - 1$ and evaluates the χ^2 of the acceptance on the round-tripped maps. The round trip is not exact on the HEALPix grid, so the stored amplitudes differ from the drawn ones. The log records the log acceptance ratio of every proposal, and at the first move of each iteration both amplitude sets are fresh: the scatter of those 3200 values is 682 ± 9 where exact draws give $\sqrt{k} = 437$, a factor 2.44 ± 0.06 in variance, and a direct calculation of the round-trip residual on the same inputs gives 2.69 ± 0.18. With the factor 2.44 in Eq. (12) the toy gives widths of 15.9 and 16.6 and an acceptance rate of 0.075 against the measured 0.074 and 0.077. The same storage residual has a gradient in the indices that predicts shifts of the chain means of $+3.1$ profile widths for the dust index and -0.1 for the synchrotron index, where $+3.0 \pm 0.5$ and -2.7 ± 2.0 were measured. The forward test, a run of the same configuration with the amplitudes stored as maps, predicted to give 13.1 and 13.6 widths and no shift, has not been made.

3.6 What this means for existing chains

A chain run with this move reports index uncertainties that are too wide by $\sqrt{T}$, a factor of order ten at map-level k, with a mean close to the profile posterior. It does not sample the flat-prior joint posterior that the rest of the Gibbs sweep assumes. The statement applies to the move as defined in steps 1 to 3 with exact conditional draws. We met it in the development version of Commander4. In Commander3 we read the code and did not run it: the index samplers used for the BeyondPlanck analysis hold the amplitudes fixed during the Metropolis step, which is a valid conditional update and not this move, while a later module (`sample_specind_mh`) proposes an index, re-solves the amplitudes by conjugate gradients and accepts on χ^2, which has the structure analysed here (`E009`). We do not know which published results, if any, were made with that module.

4 The flat amplitude prior, written in terms of the reference frequency

That the flat-prior marginal of Eq. (4) pulls spectral indices is known. Eriksen et al. (2008) found a global offset biased by more than 4σ under uniform priors and removed it with a Jeffreys prior on the index. Dunkley et al. (2009), fitting one polarized synchrotron index per $N_{\rm side} = 2$ pixel to the five-year WMAP maps with uniform priors on the CMB and synchrotron amplitudes defined at 23 GHz, found index posteriors centred on -2.6 in simulations, noted that the effect grows with the number of pixels that share an index and that a Jeffreys prior on the index does not remove it, and divided their prior by a factor estimated from a noise-only simulation. Alonso et al. (2017) cancel the determinant with a "volume prior", which in our reading makes their target the profile likelihood, and show biases of more than 3σ without it. Jeffreys or reference priors on the indices are used against prior-volume effects in several other analyses (Dickinson et al., 2009; Planck Collaboration X, 2016; Jew & Grumitt, 2020), and Andersen et al. (2023) describe the marginal and the profile ("ridge") index samplers of Commander, noting that the two differ by the determinant term. What has not been written down, as far as we can find, is how large the pull is in closed form and what sets its sign. Both follow from one identity.

4.1 An exact identity

Change the reference frequency of the power-law component from $\nu_{\rm ref}$ to $\nu'_{\rm ref}$. The mixing column of every amplitude in index region r is multiplied by $(\nu_{\rm ref}/\nu'_{\rm ref})^{\beta_r}$, so $\mathsf{M}' = \mathsf{M}\,\mathsf{D}$ with D diagonal, $\chi^2_{\rm min}$ is unchanged, and $\det \mathsf{A}' = \det \mathsf{A} \prod_r (\nu_{\rm ref}/\nu'_{\rm ref})^{2 n_r \beta_r}$, with n_r the number of power-law amplitudes in region r. Therefore

$$\ln \mathcal{L}_{\rm flat}(\beta; \nu'_{\rm ref}) = \ln \mathcal{L}_{\rm flat}(\beta; \nu_{\rm ref}) + \sum_r n_r\, \beta_r \ln \frac{\nu'_{\rm ref}}{\nu_{\rm ref}}, \tag{14}$$

for any number of components, any bandpasses and any noise covariance. Moving the frequency at which a flat amplitude prior is stated is the same as multiplying the index prior by an exponential, $e^{\lambda\beta}$ with $\lambda = n \ln(\nu'_{\rm ref}/\nu_{\rm ref})$: 34 e-folds per unit of index for 128 amplitudes moved from 23 to 30 GHz. The profile likelihood and the zero-mean prior with a free scale do not change.

4.2 The shift and its pivot

Let $g_r = -\frac{1}{2}\,\partial \ln\det \mathsf{A}/\partial\beta_r$ be the gradient of the volume term. To first order the maximum of $\mathcal{L}_{\rm flat}$ sits at the profile maximum plus

$$\delta\beta = \mathsf{H}^{-1} g, \qquad \mathsf{H} = \mathsf{J} + \nabla\nabla\, \tfrac{1}{2} \ln\det \mathsf{A}, \tag{15}$$

where H is the expected curvature of the flat-prior marginal itself: the profile information plus the curvature of the volume term. For one power-law component with delta bands and diagonal noise, $-\frac{1}{2}\ln A_p = \beta \ln \nu_{\rm ref} - \frac{1}{2}\ln \sum_b w_{pb}\, \nu_b^{2\beta}$, so

$$g = \sum_p \ln \frac{\nu_{\rm ref}}{\bar\nu_p}, \qquad \ln \bar\nu_p = \frac{\sum_b w_{pb}\, \nu_b^{2\beta} \ln \nu_b}{\sum_b w_{pb}\, \nu_b^{2\beta}}, \qquad \nabla\nabla\, \tfrac{1}{2} \ln\det \mathsf{A} = 2 \sum_p V_p = 2\,\mathsf{T}, \tag{16}$$

where $\bar\nu_p$ is the frequency at which amplitude p is best measured: the geometric mean of the band frequencies under the weights (inverse noise variance times squared spectrum) that also define the lever arm V_p in Eq. (7). When the amplitudes of a region share a noise pattern, $\bar\nu_p = \nu_{\rm piv}$ and, with $\sigma_\beta^2 = 1/\mathsf{J}$ the profile variance of the region index,

$$\boxed{\delta\beta = \ln\!\left(\frac{\nu_{\rm ref}}{\nu_{\rm piv}}\right) \frac{n\,\sigma_\beta^2}{1 + 2\,\mathsf{T}/\mathsf{J}} \;\longrightarrow\; \ln\!\left(\frac{\nu_{\rm ref}}{\nu_{\rm piv}}\right) \sigma_1^2 \quad (\mathsf{T} \ll \mathsf{J})} \tag{17}$$

where $\sigma_1^2 = n\sigma_\beta^2$ is the index variance that a single amplitude carries, and T/J is the inverse mean squared signal-to-noise ratio of the amplitudes (Sect. 5). The boxed form is for one component. With further free amplitudes the curvature of the volume term is no longer $2\mathsf{T}$ and can be negative, so that H is smaller than J and the flat-prior marginal may have no interior optimum at all; Sect. 6.6 shows a case. Three things follow.

The shift does not average down. Adding amplitudes at fixed depth leaves σ_1 and T/J unchanged, so $\delta\beta$ stays the same while the posterior width falls as $n^{-1/2}$. In units of the posterior width the pull grows as $\sqrt{n}$: this is the observation of Dunkley et al. (2009) that the effect increases with the number of pixels per index, and it is the Neyman-Scott structure (Neyman & Scott, 1948). More data of the same kind make it worse. So does a finer pixelization of the same data: for $\mathsf{T} \ll \mathsf{J}$, four times as many amplitudes for the same region information is four times the shift.

The sign is set by the convention. The index is pulled steep if the reference lies below the pivot and flat if above. A flat prior is neutral only when stated at the pivot.

The pivot depends on the other components and on their priors. A component with a free amplitude and a spectrum that does not depend on the index (the CMB, dust with fixed parameters) contributes nothing to g directly, but projecting it out replaces the weights w_{pb} by the projected weight matrix $\tilde{\mathsf{W}}_p = \mathsf{W}_p - \mathsf{W}_p\mathsf{F}(\mathsf{F}^\mathsf{T}\mathsf{W}_p\mathsf{F} + \mathsf{S}^{-1})^{-1}\mathsf{F}^\mathsf{T}\mathsf{W}_p$, with F the spectra of the other components and S their prior covariance. The pivot becomes a generalized mean under $\tilde{\mathsf{W}}_p$, which has negative off-diagonal elements and can lie outside the range of the bands. With a free CMB amplitude under a flat prior the index information sits in the lowest band and the pivot moves to or below it. With a CMB prior the size of the signal it moves back.

For several regions, several components and dense noise, g and H are a vector and a matrix, and Eq. (14) gives the difference between two references,

$$\hat{\beta}_\mathrm{flat}(\nu'_\mathrm{ref}) - \hat{\beta}_\mathrm{flat}(\nu_\mathrm{ref}) \simeq \mathsf{H}^{-1}\left[\, n_r \ln(\nu'_\mathrm{ref}/\nu_\mathrm{ref}) \,\right], \tag{18}$$

in which the regions are coupled: a large region with many amplitudes drags its neighbours through the off-diagonal elements of H^{-1}. We test this form on simulations and on real data in Sect. 6. Equation (15) was corrected by an independent review of our first derivation, which had J in place of H; the two agree where $\mathsf{T} \ll \mathsf{J}$, and a direct test with the reference placed close to the pivot gives a measured shift of 1.01 and 1.03 times Eq. (15) at $\mathsf{T}/\mathsf{J} = 0.03$ and 0.11 where the uncorrected form is off by 5 and 16 percent (E039, independent review). For shifts that are not small against the range over which the curvature changes, neither is accurate (Fig. 2, right).

4.3 Sizes at the depth of the public polarization data

Table 2 gives the sizes on a realistic polarized sky at the white-noise depth of the WMAP and LFI maps. For the synchrotron-only model the pivot of the five lowest bands is 26.4 GHz, and references of 23 and 30 GHz, the two in use, sit on opposite sides of it: the same data give an index steeper than the truth by 0.06 in one convention and flatter by 0.06 in the other, one posterior standard deviation each way, for regions measured to 0.06. For regions measured to 0.16 the pull is 0.4 and 0.3. The difference between the two references follows Eq. (14): $\ln(30/23) \times 128 \times 0.059^2 = 0.118$ against a measured 0.123. Doubling and quadrupling the noise leaves the shift at a given σ_β unchanged, as Eq. (17) requires (-0.063 and -0.073 at 23 GHz for the class 0.03 to 0.1; E036). The profile likelihood on the same simulations is unbiased to 0.002 in the first class and 0.02 in the second (Sect. 5).

With free CMB and dust amplitudes under flat priors the pivot moves to 22 to 23 GHz. A flat prior stated at 23 GHz is then nearly neutral and one stated at 30 GHz pulls the index flat by 0.17 to 0.86. With a Gaussian CMB prior of 0.2 μK per element, the size of the CMB polarization signal at this resolution, the pivot and the shifts return to the synchrotron-only values, and so does the index error, which the flat CMB prior had inflated by a factor 1.63. The prior on

Table 2: The flat-prior marginal against Eq. (17) on simulations at the white-noise depth of the WMAP and LFI polarization maps (E036). One index per $N_{\rm side} = 2$ pixel (128 Q/U amplitudes at $N_{\rm side} = 16$), exact one-dimensional posteriors on a grid, 100 noise realizations; truth: the Cosmoglobe DR1 polarized synchrotron map smoothed to 5° with $\beta = -3.1$. Regions are grouped by their profile index error σ_β. "Shift" is the mean of posterior mean minus truth. The predicted value is the high signal-to-noise limit of Eq. (17) with the class median of σ_β (in brackets) and the pivot of the band set; it is first order in the shift and uses one number for a class, so it degrades where the shift exceeds the posterior width.

model	σ_β class (median)	reference	shift	predicted	shift / posterior sd
synchrotron only,	below 0.03 (median 0.019)	23 GHz	−0.005	−0.006	−0.26
bands K, 30, Ka, Q, 44		30 GHz	+0.007	+0.006	+0.35
($\nu_{\rm piv} = 26.4$ GHz)		60 GHz	+0.037	+0.038	+1.9
	0.03 to 0.1 (0.059)	23 GHz	−0.061	−0.062	−0.94
		30 GHz	+0.062	+0.057	+1.01
		60 GHz	+0.341	+0.37	+6.1
	0.1 to 0.25 (0.161)	23 GHz	−0.40	−0.46	−2.4
		30 GHz	+0.34	+0.42	+2.5
free CMB and dust, flat priors,	0.03 to 0.1	23 GHz	+0.006		
adding V, 70, 353 GHz		30 GHz	+0.167		
	0.1 to 0.25	23 GHz	+0.058		
		30 GHz	+0.859		
the same with a CMB prior	0.03 to 0.1	23 GHz	−0.060		
of 0.2 μK per element		30 GHz	+0.063		

the CMB amplitude matters more for the synchrotron index than the prior on the synchrotron amplitude does, as long as the latter is not flat.

An existing real-data comparison. An earlier experiment of ours (E022) evaluated both the flat-prior marginal and the profile likelihood on real WMAP Ka to V, LFI and 353 GHz polarization maps in the regions of Svalheim et al. (2023), at a 30 GHz reference, for two noise models. In its high signal-to-noise limit Eq. (17) says the difference of the two estimates divided by $n\sigma_\beta^2$ is $\ln(30/\nu_{\rm piv})$ whatever the noise model. In the Spur region the two noise models give shifts of +0.271 and +0.172 for $n\sigma_\beta^2 = 1.913$ and 1.216: a pivot of 26.04 GHz in both cases, to 0.2 percent, for a band set whose lowest band is at 28.4 GHz and with a free CMB amplitude. With a CMB prior of 0.19 μK the implied pivot is 29.5 to 30.2 GHz. For that band set the neutral reference frequency is therefore 26 GHz or 30 GHz depending on how the CMB is constrained. For WMAP K to V on the real maps it is 23.3 GHz with a CMB prior of the size of the signal and about 20 GHz, below the lowest band, with a flat one (Sect. 6). Sect. 6 tests the identity directly on the WMAP maps with their dense covariances.

5 Two targets without the volume term, and an honest error for the simpler one

5.1 The profile likelihood: unbiased score, curvature too sharp

For amplitudes that enter linearly with known noise, the profile score is exactly unbiased whatever the true amplitudes are (App. A). The amplitudes are fixed unknowns, no distribution is assumed for them, and nothing depends on the reference frequency. This is the functional errors-in-variables model with known error covariance. There the consistency of the slope is classical, and so is the caveat that the variance of the estimator is not the one derived from the Fisher information matrix (Engel, 1985, and references therein); Stompor et al. (2009) remark of the profile curvature that it gives "useful, if not strictly precise" uncertainties. What follows is the

form of that caveat for several components and dense noise, with its size. In the notation of Eq. (6) the score variance is $\mathsf{J}+\mathsf{T}$ while the curvature of $\frac{1}{2}\chi^2_{\min}$ is J, so the estimator has covariance

$$\mathrm{Cov}(\hat{\beta}) = \mathsf{J}^{-1}(\mathsf{J}+\mathsf{T})\,\mathsf{J}^{-1}, \tag{19}$$

a sandwich (Huber, 1967; White, 1982) whose middle term is known in closed form here. An interval built from the curvature of the profile likelihood, or from an MCMC chain run on Eq. (5), is too narrow by

$$\frac{\mathrm{sd}(\hat{\beta})}{\sigma_{\mathrm{curv}}} = \sqrt{1+\frac{\mathsf{T}}{\mathsf{J}}} = \sqrt{1+\frac{\sum_p V_p}{\sum_p u_p V_p}}, \tag{20}$$

the second form for one region and one component. The ratio T/J is the inverse of the lever-arm-weighted mean squared signal-to-noise ratio of the amplitudes. For a region measured through a few bright pixels it is negligible. For an index shared by many faint pixels it is not, however well the index itself is measured: this is why the profile likelihood was found unbiased but "up to two times too narrow" for a full-sky index at low signal to noise in our earlier simulations (E017), and why it is safe for the bright regions of the polarized sky (E036).

5.2 The error from one data set

On real data J has to be evaluated at fitted amplitudes, and here is a trap. The fitted amplitudes are the true ones plus noise with covariance A^{-1}, and

$$\mathrm{E}\big[\mathsf{J}(\hat{a})\big] = \mathsf{J}+\mathsf{T}\,: \tag{21}$$

the plug-in information estimates the score variance, not the curvature. A sandwich built as $\mathsf{J}(\hat{a})^{-1}[\mathsf{J}(\hat{a})+\mathsf{T}]\mathsf{J}(\hat{a})^{-1}$ is therefore narrower than the curvature width it was meant to widen. The consistent estimate is

$$\mathsf{J}_{\mathrm{c}} = \mathsf{J}(\hat{a}) - \mathsf{T}, \qquad \widehat{\mathrm{Cov}}(\hat{\beta}) = \mathsf{J}_{\mathrm{c}}^{-1}\,\mathsf{J}(\hat{a})\,\mathsf{J}_{\mathrm{c}}^{-1}, \tag{22}$$

and a region for which J_{c} is not positive has fitted amplitudes consistent with noise and no measured index. Both $\mathsf{J}(\hat{a})$ and T come from quantities already formed in the fit.

Figure 2 tests this over a ladder designed to reach the regime where it matters: many faint amplitudes sharing one index, with predictions written down before the run (E041). With the curvature width the standard deviation of the standardized error equals Eq. (20) within 0.054 in all 30 cells where the index is measured, up to 1.59 at $\mathsf{T}/\mathsf{J}=1.56$, and 68 percent intervals cover the truth as rarely as 49 percent of the time. With the plug-in sandwich they cover it 38 percent of the time. With Eq. (22) the standardized error has standard deviation 0.95 to 1.05 and the coverage is 0.65 to 0.72 in 29 cells and 0.635 in one; the mean standardized error is within 0.10 everywhere. Sect. 6 repeats the test of the error taken from each realization for several coupled regions, three components and the dense WMAP covariances, in a regime where T/J is small. Two limits of the recipe, found by an independent check: the first-order sandwich is about 3 percent narrow on average and up to 7 percent at $\mathsf{T}/\mathsf{J}\simeq 1$, and above $\mathsf{T}/\mathsf{J}\simeq 1.5$ the plug-in estimate is too noisy to use even when J_{c} is positive, so we quote an error only where $\mathsf{T}/\mathsf{J}_{\mathrm{c}}<1$. With a Gaussian prior on a component that does not depend on the index, such as the CMB, the score is unbiased over skies drawn from that prior and not for a sky whose variance differs from it; the first-order bias is given in App. A.

5.3 The zero-mean prior with a free scale

The Bayesian relative of the profile likelihood is a zero-mean Gaussian amplitude prior with a variance per index region that is inferred with the index. For a noise pattern that is common

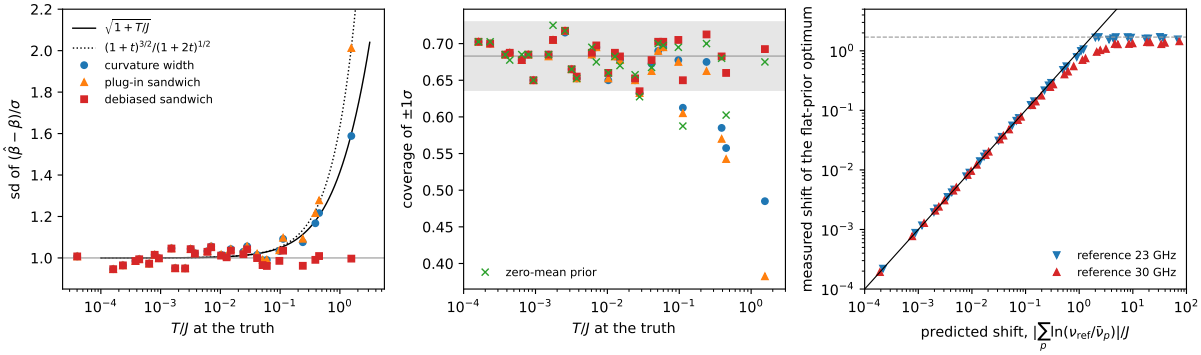


Figure 2: One index shared by $n = 32$ to 8192 heavy-tailed amplitudes at five bands, per-amplitude signal-to-noise ratio from 9 to 0.07, 400 noise realizations per cell; cells in which the index is measured to better than 0.25 (E041). Left: standard deviation of the standardized error of the profile optimum with three error estimates, each taken from the same data set as the estimate, against T/J at the truth. The curvature width follows Eq. (20) (line), the plug-in sandwich is narrower still (dotted: its expected factor), and the debiased sandwich of Eq. (22) stays at one. Centre: coverage of the $\pm1\sigma$ interval for the same three and for the zero-mean prior; the band is ±2 binomial standard errors around 0.683. Right: shift of the flat-prior optimum from the profile optimum against g/J, the high signal-to-noise limit of Eq. (15), at references of 23 and 30 GHz, all 40 cells; the dashed line is the edge of the allowed range.

to the amplitudes of a region, $w_{pb} = w_b$, the marginal likelihood depends on the data through $q(\beta) = \sum_p b_p^2/A$ only, with $b_p = \sum_b w_b m_b d_{pb}$ and $A = \sum_b w_b m_b^2$, and maximizing over τ gives

$$\ln \mathcal{L}_\tau(\beta) = \tfrac{1}{2}\, q(\beta) - \tfrac{n}{2} \ln q(\beta) + \text{const} \qquad (q > n), \tag{23}$$

a monotonic function of the profile log-likelihood $\frac{1}{2}q$. The optimum is the profile optimum, and the curvature is the profile's times $1 - n/q$, which is the plug-in form of $1/(1 + \mathsf{T}/\mathsf{J})$. In this case (one component, a common noise pattern, the scale maximized) the zero-mean prior with a free scale *is* the profile likelihood with its sandwich width: the prior learns from the scatter of the fitted amplitudes how much of them is noise. Simulation confirms the identity to the tolerance of the optimizer (E041).

When the noise pattern varies between the amplitudes, as it does on real maps, the two estimators part. The Gaussian prior is then a model for the amplitudes, real polarized amplitudes are far from Gaussian, and the bright ones carry whatever noise pattern they happen to have. In the ladder the resulting bias of the zero-mean prior is between -0.6 and $+0.2$ posterior standard deviations in the cells where the index is measured, with a sign that belongs to the amplitude configuration and not to the noise level, while its coverage stays within 0.59 to 0.73. On the realistic sky of Table 2 it was within 0.11 standard deviations except in one class at four times the real noise (-0.30; E036). A scale map in place of a single scale reduces it there. We therefore take the profile likelihood with the debiased sandwich as the estimator of record and report the zero-mean prior beside it. What the prior offers in exchange is a generative model for the amplitudes, which a hierarchical model of an index field needs.

5.4 What each target is good for

Table 3 collects the results. Two remarks on scope. First, for an index *field* with a hierarchical prior the profile likelihood is not enough at low signal to noise: the profile maxima of faint regions scatter, and a hierarchy reads that scatter as field variance (E012). A proper amplitude prior is then needed, and its specification becomes part of the model (Sect. 7). Second, the statements above are for amplitudes that enter linearly with known Gaussian noise. Beams, masks and a correlated-noise model that is itself uncertain change the numbers, not the structure.

Table 3: The index targets of Sect. 2 for an index shared by many amplitudes, as established by the experiments cited.

target	centre	width	depends on ν_{ref}
flat-prior marginal (and Gibbs with flat amplitude priors)	shifted by $\ln(\nu_{\mathrm{ref}}/\nu_{\mathrm{piv}})\,\sigma_1^2$ (E036, E041)	curvature	yes: tilt $e^{n\beta \ln \nu_{\mathrm{ref}}}$
profile likelihood	unbiased (E036, E041)	curvature too narrow by $\sqrt{1+\mathsf{T}/\mathsf{J}}$; debiased sandwich calibrated (E041)	no
zero-mean prior, free scale per region	unbiased for a common noise pattern; up to 0.6 sd otherwise (E041)	calibrated where unbiased	no
re-solve move inside a Gibbs sweep	close to the profile optimum	too wide by $\sqrt{T}$, Eq. (10) (E035)	no

6 Demonstration on the public polarization maps

This section does one thing: it applies the estimators of Sects. 4 and 5 to the same real maps, with the same regions, bands and noise model, so that the difference between their answers is the difference between the estimators and nothing else. It does not deliver a polarized synchrotron index. Sect. 6.4 shows why not: at the statistical precision these maps offer, several analysis choices that are independent of the estimator each move the index by more than its error.

6.1 Data and model

WMAP. The nine-year low-resolution ($N_{\mathrm{side}} = 16$) Stokes Q and U maps of the K, Ka, Q and V bands (Bennett et al., 2013) with their released dense inverse noise covariance matrices, which carry the time-correlated noise of the differential scan and de-weight the poorly constrained transmission-imbalance modes (Jarosik et al., 2007, 2011). We checked them as released (E038): symmetric to 10^{-9}, positive definite with no null space, a pixel noise of 2.0 to 2.2 μK in brightness temperature, and a few modes at multipoles 2 to 6 whose variance is 65 to 164 times the median. Watts et al. (2023) argue from a joint reprocessing that these matrices do not fully describe the large-scale uncertainty of the nine-year maps; we use them as released and test the sensitivity to the modes they flag.

Planck. The PR4 maps (Planck Collaboration Int. LVII, 2020): 353 GHz as the dust tracer, and 30 GHz in Sect. 6.5. For 353 GHz the per-pixel Q/U covariance of the map file is combined over the sub-pixels of each $N_{\mathrm{side}} = 16$ pixel, and the map is the corresponding weighted average.

Spectra. Each band enters through its measured bandpass. For WMAP every detector is given unit response to a thermal spectrum and the detectors and radiometers of a band are averaged with equal weights, which reproduces the per-radiometer effective frequencies of the WMAP team's routine to 0.005 GHz, and the team's correction for the drift of the band centres over the mission (0.13 percent at K band) is applied. For a power law of index -3.1 the effective frequencies are 22.41, 32.66, 40.30 and 59.99 GHz for K to V and 28.09 GHz for the 30 GHz band. Delta bands at the nominal WMAP frequencies would steepen a K to Ka index by about 0.06 (a calculation, not a fit). The 353 GHz band cannot be a delta band: a modified blackbody with index 1.53 and temperature 19.6 K (Planck Collaboration XI, 2020) seen through the band is 1.12 to 1.14 times its value at 353.0 GHz, depending on which bolometers are counted, so a delta band overstates the dust at every lower frequency by that factor. We use the band of the polarization-sensitive bolometers.

Beams. The low-resolution maps are pixel averages of maps with beams from 0.88° (K) to 0.08° (353 GHz). A noise-free simulation with index −3.1 everywhere, each band seen through its own Gaussian beam, gives region indices biased flat by 0.005 to 0.011 for the WMAP bands (E038). We therefore bring every band to a common 1° window with an invertible operator on the $N_{\mathrm{side}} = 16$ grid, $\mathsf{S}_b = 1 + \mathsf{Y}\,\mathrm{diag}(f_\ell - 1)\,\mathsf{Y}^+$ with f_ℓ the ratio of the beam windows to multipole 47, transforming the data and the covariance together. In the simulation the residual bias is 0.002 or less, and on the real maps the correction moves the four indices by 0.0068, 0.0025, 0.0085 and 0.0035 where the simulation predicted 0.0070, 0.0030, 0.0088 and 0.0041 (white CMB prior).

Model. Synchrotron as a power law with one index per region and one amplitude per element (pixel and Stokes parameter); dust as the modified blackbody above with one amplitude per element under a flat prior; and the CMB with one amplitude per element. The prior on the CMB amplitudes is the ΛCDM covariance of E and B modes on the grid: spectra from CAMB (Lewis et al., 2000) for the Planck 2018 parameters (Planck Collaboration VI, 2020) with an optical depth of 0.054, through the working window and the pixel window, for multipoles 2 to 47, the span of the harmonic basis of the grid, plus a white floor of 0.01 μK. The pixel average of a ΛCDM sky has a further 0.11 μK per element from higher multipoles, nearly white, which this prior does not describe (independent review, E040); with the floor raised to that value the four indices change by 0.0006 or less and their errors by 0.3 percent. Two other choices are shown beside it, because the choice matters both for the index and for the pivot of Sect. 4: a white Gaussian prior of 0.2 μK per element, which we used first, and a flat prior, for which the profile score is unbiased whatever the CMB is. The white prior has nearly the right total variance (0.17 μK per element for ΛCDM) and the wrong shape: it understates the ΛCDM E-mode variance by a factor of 200 at the quadrupole and overstates it above multipole 10. Regions: the four of Svalheim et al. (2023) (High Latitude, Spur, Galactic Center, Plane; 4388, 928, 182 and 646 elements). No mask is applied.

Expectations. The expectations for the baseline fit and its estimator comparison were written down, with numbers, before the first fit to real data, and the simulations of the next subsection were read first (E039, E040). The fits of Sects. 6.4 to 6.6 were each preceded by written expectations too, but those were written after earlier real-data results had been seen; the records mark which is which, and record the expectations that failed.

6.2 Validation on simulations with the real covariances

Before any fit to real data the chain was run on simulations with known indices: synchrotron amplitudes from the Cosmoglobe DR1 polarized synchrotron map, dust amplitudes from the 353 GHz map, a CMB drawn from the white prior, one true index per region chosen away from the centre of the allowed range (−3.15, −3.25, −2.85, −2.95), and noise drawn from the real covariances (E039). Six criteria were fixed in advance.

Table 4 gives the outcome for the four regions at $N_{\mathrm{side}} = 8$, where the simulated sky is constant inside each pixel by construction. The profile optimum is unbiased to 0.07 standard deviations, its sandwich error is right to 6 percent, the 68 percent intervals cover the truth 69 percent of the time, and the reduced χ^2 at the optimum is 1.0006 with a scatter of 0.022. A spot check at the working resolution, $N_{\mathrm{side}} = 16$, with 12 realizations gives a pooled standardized error of mean +0.03 and standard deviation 0.99 and a reduced χ^2 of 0.998 ± 0.009. The zero-mean prior lands within 0.06 standard deviations of the profile optimum. For this band set with the white CMB prior the pivot sits at 23 GHz: a flat prior stated there is neutral, and the same prior stated at 30 GHz moves the High Latitude index by 4.5 standard deviations. The single-region formula predicts that move for the large region and fails for the three small ones, which move by 1.7 to 5 times their own prediction: they are dragged by the large region through the shared noise. The coupled form, Eq. (18), is within 5 percent in all four. Two of our six criteria failed as written,

Table 4: Validation of the region-index estimators on simulations with the released WMAP noise covariances (E039): 200 noise realizations at $N_{\mathrm{side}} = 8$, the first 40 with every estimator. z is the error of the profile optimum in units of the sandwich error at the truth. The last block compares the difference between the flat-prior optima at 30 and 23 GHz with Eq. (17) applied to each region alone and with the coupled form, Eq. (18).

	High Latitude	Spur	Galactic Center	Plane
elements	1124	212	50	150
sandwich error at the truth	0.0153	0.0166	0.0108	0.0146
mean of z	-0.02 ± 0.07	0.00 ± 0.07	$+0.07 \pm 0.07$	$+0.01 \pm 0.07$
standard deviation of z	1.04	0.98	0.94	1.02
coverage of $\pm 1\sigma$	0.670	0.715	0.685	0.695
zero-mean prior minus profile, in σ	+0.03	0.00	+0.06	+0.04
flat prior at 23 GHz minus profile	−0.0008	0.0000	−0.0002	−0.0004
flat prior at 30 GHz minus profile	+0.0680	+0.0270	+0.0074	+0.0138
in units of σ	+4.5	+1.6	+0.7	+0.9
30 GHz minus 23 GHz, measured	+0.0687	+0.0270	+0.0076	+0.0142
predicted, each region alone	+0.0692	+0.0156	+0.0015	+0.0085
predicted, coupled	+0.0723	+0.0278	+0.0078	+0.0145

one because it was worded wrongly (a lower bound of 1.0 on a ratio that scatters around one) and one because it was written for the single-region formula. In these regions T/J is below 0.02, so the errors taken from each realization (Eq. 22) differ from the errors at the truth by less than a percent: with the errors taken from each of 40 realizations the standardized error has a standard deviation of 0.97 and the 68 percent coverage is 0.70, for the curvature, plug-in and debiased errors alike. The correction itself is tested in Fig. 2, not here. These simulations were first run on a numerical path that was later found to overwrite memory (Sect. 8); they were run again on a sound path and every number in Table 4 and in this paragraph was reproduced to the last printed digit (E043).

6.3 The same maps, several answers

Table 5 is the result of this section. The numbers in each column come from the same maps, the same regions, the same bands, the same bandpasses and the same noise covariances. They differ only in how the amplitudes are removed.

Under flat amplitude priors the answer depends on where the amplitudes are defined. For this band set with a ΛCDM prior on the CMB the pivot lies just above 23 GHz (23.3 to 23.4 GHz for the four regions, to first order), and a flat prior stated at 23 GHz agrees with the profile optimum to 0.8 standard deviations. The same prior stated at 30 GHz gives a High Latitude index of −2.92 instead of −3.08: flatter by 0.16, 12.6 standard deviations of the sandwich error, with no change in the data or the model. The Spur moves by 0.09 and the Plane by 0.05. The difference between the two conventions is 0.91 to 0.95 of the first-order prediction, Eq. (18), in all four regions, so the identity of Sect. 4 holds on real maps at the predicted size. The zero-mean prior with a free scale is within 0.8 standard deviations of the profile optimum in every region.

The prior on the CMB. The CMB has one amplitude per element like the other components, and the index responds to what that amplitude is allowed to be (Table 6). We began with a white prior of 0.2 μK per element and had fixed in advance that the ΛCDM covariance would replace it if any index moved by more than one standard deviation between the two. The Plane moved by 2.3 and the Spur by 1.0: with the white prior all four indices are steeper, by 0.012, 0.017, 0.006 and 0.031. A simulation says what to make of that (E040): 60 skies with the synchrotron, dust and noise of the validation above and a CMB drawn from the ΛCDM covariance, each fitted with the three priors at $N_{\mathrm{side}} = 8$. With the matched prior the estimate is unbiased and its error is

Table 5: The polarized synchrotron index of the WMAP nine-year K, Ka, Q and V maps (with Planck PR4 353 GHz for dust) in the four regions of Svalheim et al. (2023), at $N_{\rm side} = 16$ with the dense noise covariances, all bands brought to a common window, and the ΛCDM covariance as the prior on the CMB (E040). Top: the profile optimum with two errors, the debiased sandwich error, which assumes one index per region (in parentheses for the Galactic Center, where that model does not fit), and a block error from the scatter of the score over patches of sky, which does not (see the text; none is quoted for the Galactic Center, where fewer than three patches carry the information). Then the answer when the same likelihood is instead marginalized over flat amplitude priors with the synchrotron amplitude defined at 23 or at 30 GHz, or under a zero-mean prior with a free scale. Middle: the difference between the two flat-prior answers against Eq. (18). Bottom: goodness of fit, as χ^2 per element and band, for which three free amplitudes on five bands leave an expectation of 0.6.

	High Latitude	Spur	Galactic Center	Plane
elements	4388	928	182	646
profile likelihood	−3.082	−3.080	−3.081	−2.964
sandwich error	0.012	0.016	(0.009)	0.013
block error	0.044	0.038		0.100
effective number of patches	5.7	8.8	2.5	5.1
flat prior at 23 GHz	−3.092	−3.085	−3.082	−2.967
flat prior at 30 GHz	−2.925	−2.990	−3.054	−2.919
minus profile, in units of the sandwich error	+12.6	+5.4	+3.1	+3.4
zero-mean prior, free scale	−3.090	−3.093	−3.075	−2.972
flat 30 GHz minus flat 23 GHz	+0.167	+0.095	+0.029	+0.048
Eq. (18)	+0.184	+0.102	+0.030	+0.051
χ^2 per element and band	0.58	0.59	1.78	0.77
T/J_c	0.014	0.005	0.0003	0.003

Table 6: The fit of Table 5 under three priors on the CMB amplitudes: the ΛCDM covariance, a white prior of 0.2 μK per element, and a flat prior (E040).

	High Latitude	Spur	Galactic Center	Plane
profile likelihood, with the sandwich error				
ΛCDM covariance	-3.082 ± 0.012	-3.080 ± 0.016	-3.081 ± 0.009	-2.964 ± 0.013
white, 0.2 μK	-3.094 ± 0.012	-3.096 ± 0.015	-3.087 ± 0.009	-2.995 ± 0.012
flat	-3.028 ± 0.029	-3.055 ± 0.036	-3.008 ± 0.020	-2.988 ± 0.030
flat amplitude prior at 23 GHz, minus profile				
ΛCDM covariance	−0.010	−0.005	−0.002	−0.003
white, 0.2 μK	−0.008	−0.003	−0.001	−0.002
flat	+0.517	+0.235	+0.074	+0.124
flat amplitude prior at 30 GHz minus at 23 GHz: measured (Eq. 18)				
ΛCDM covariance	+0.167 (+0.184)	+0.095 (+0.102)	+0.029 (+0.030)	+0.048 (+0.051)
white, 0.2 μK	+0.158 (+0.172)	+0.075 (+0.079)	+0.025 (+0.027)	+0.037 (+0.038)
flat	+0.799 (+0.874)	+0.380 (+0.423)	+0.136 (+0.136)	+0.208 (+0.225)
χ^2 per element and band (expectation 0.6, 0.6 and 0.4)				
ΛCDM covariance	0.584	0.592	1.782	0.770
white, 0.2 μK	0.581	0.591	1.815	0.776
flat	0.382	0.395	0.643	0.404

calibrated (standardized errors with mean −0.15 to +0.11 and standard deviation 0.93 to 1.14 by region). With the white prior it is unbiased on average, but the quoted error is too small by 24 to 32 percent in three of the four regions, and the 68 percent intervals cover 48 to 65 percent of the time: the large-scale E modes that a white prior does not allow move the estimate from sky to sky, by 0.007, 0.009, 0.003 and 0.007 relative to the matched fit. With a flat prior the estimate is unbiased and calibrated at 2.2 to 2.5 times the error. Measured against that sky-to-sky scatter and its correlation between regions, the moves seen on the real maps at High Latitude, in the Spur and in the Galactic Center are unremarkable (22 percent of the simulated skies lie further out), and the move in the Plane is 4.4 standard deviations, which a ΛCDM sky does not do. [The same scatter at $N_{\rm side} = 16$.] A white prior of the right total variance is therefore not a harmless stand-in: it leaves the index unbiased and makes its error bar too small by a quarter to a third before any misfit.

With a flat prior on the CMB the profile score is unbiased whatever the CMB is, at the price of errors 2.2 to 2.3 times larger, and the indices are -3.03 ± 0.03, -3.06 ± 0.04, -3.01 ± 0.02 and -2.99 ± 0.03. Against the ΛCDM version, three regions are within the scatter that the simulation gives for that difference (1.8, 0.7 and −0.7 of it) and the Galactic Center is 3.4 away. The two regions whose index depends on the CMB prior by more than a ΛCDM sky allows are the two that do not fit.

A free CMB amplitude also moves the pivot of Sect. 4, and with it the size of the convention effect. With a flat CMB prior and flat priors on the other amplitudes, the High Latitude optimum of these same four WMAP bands is at −2.51 when the synchrotron amplitude is defined at 23 GHz and at −1.71 when it is defined at 30 GHz, where the profile likelihood says -3.03 ± 0.03: the pivot has moved below the K band (to 19.7 to 19.9 GHz, to first order), and both conventions are far from it. The simulated band set of Table 2, which includes the LFI bands, has its pivots about 3 GHz higher in both cases and shows the same drop when the CMB is freed. The difference between them, 0.80, is 0.92 of the first-order prediction, and the prediction is within 10 percent in every region and for every prior in Table 6, for shifts between 0.03 and 0.8.

Errors that do not assume the model. The sandwich error is the error of the estimator if each region has one index and the noise model is right. The χ^2 per element and band says that this is close to true at High Latitude and in the Spur (0.58 and 0.59 against 0.6), 30 percent off in the Plane and wrong by a factor of three in the Galactic Center. An error that does not lean on it comes from the data: at the optimum the score of a region is a sum over elements, $\sum_p g_p = 0$, and under the model the covariance of that sum is $\mathsf{J} + \mathsf{T}$; if the index varies inside the region the partial sums over patches of sky scatter by more. With the 48 pixels of an $N_{\rm side} = 2$ grid as patches, $G_B = \sum_{p \in B} g_p$, the block covariance is

$$\widehat{\mathrm{Cov}}_{\rm block}(\hat\beta) = \mathsf{J}_{\rm c}^{-1} \Big[\sum_B \frac{G_B G_B^{\mathsf{T}}}{(1-h_B)^2} \Big] \mathsf{J}_{\rm c}^{-1}, \qquad h_B = \frac{D_B}{\sum_{B'} D_{B'}}, \tag{24}$$

with $D_B = \partial G_B / \partial \beta$ the information carried by patch B. Without the factor $(1-h_B)^{-2}$ this is the cluster sandwich in common use (Liang & Zeger, 1986), which is biased low when a few clusters carry the information; with it, it is the leverage-corrected form that in linear regression equals the delete-one-cluster jackknife variance up to a factor $G/(G-1)$ for G clusters (Bell & McCaffrey, 2002; Cameron & Miller, 2015). Our estimator is not linear, so the correspondence holds to first order, and it is exact only where neither the noise covariance nor the prior on the CMB couples patches. Patches are treated as independent. The information is far from evenly spread: the number of patches that would carry it if all carried the same, $(\sum_B D_B)^2 / \sum_B D_B^2$ (an analogue, in shares of the information, of the effective number of clusters of Carter et al. 2017), is 6, 9, 3 and 5 for the four regions on the maps, out of 46, 29, 4 and 16 patches with data. A block error from one sky is therefore itself uncertain by a third or so and should be read with the quantiles of a t distribution with few degrees of freedom (Bell & McCaffrey, 2002; Cameron & Miller, 2015); for the Galactic Center we quote none.

What a block error means needs care (independent review, E040). It is the error of the estimate as a measurement of the mean index of a population of exchangeable patches. For the information-weighted mean index of the region as realized on this sky, the sandwich error is the right one. If the index varies coherently over several patches the block error is too small (0.6 to 0.7 of the true scatter for a correlation length of two patches in a test by the reviewer). And if the noise covariance or the prior couples patches, the patch scores are anticorrelated and the block error is too large even when the model is true, so the ratio of block error to sandwich error under the model has to be taken from simulations with the covariances and the prior of the fit. With independent patches (white CMB prior, $N_{\rm side} = 8$) the block error reproduces the scatter of the estimate: the root mean square of the standardized error is 0.82, 1.20 and 1.02 at High Latitude, in the Spur and in the Plane where 1.1 to 1.2 is expected, and the form without the leverage correction is too small (1.39 and 1.26 in the Spur and the Plane). When the same simulations are given a true spread of the index of 0.1 between patches, drawn anew for each realization, the sandwich error is too small by a factor of 2.4 to 3.1 and the block error follows the scatter of the estimate about the mean index of the region (0.59, 0.99 and 0.91). In that test the scatter of the estimate is, within 10 percent, the spread between patches divided by the square root of the effective number of patches, added in quadrature to the sandwich error. The rule for choosing between this form and two others, closest to one in both tests and in doubt the more conservative, was written before either test was read; the three forms differ by 15 to 45 percent on the maps and all are in the record.

On the maps the block error is 3.5, 2.3 and 7.5 times the sandwich error at High Latitude, in the Spur and in the Plane (Table 5, Fig. 3): 0.04, 0.04 and 0.10 where the model gives 0.012, 0.016 and 0.013. Under the model, with the window-matched covariances and the ΛCDM prior of the fit, the ratio is not one: in 60 simulations at $N_{\rm side} = 8$ its median is 1.28, 1.21 and 1.32 and its largest value 2.8, 2.4 and 2.8, because that prior couples the patches (with the white prior the medians are 1.25, 1.03 and 1.07). Against that reference the maps are outside the range at High Latitude and in the Plane, and the Spur is inside it: one simulation of 60 has a larger ratio. [The same reference at $N_{\rm side} = 16$.] So the High Latitude region and the Plane are not one index each; for the Spur the block error is an upper limit on the error and no evidence of misfit. In the simulation with a spread of 0.1 the median ratio is 3.8, 2.4 and 3.6. (That simulation has the white prior; a spread of that size is the scale of what the High Latitude ratio asks for, and the Plane asks for more.) What the fit of a region returns is the index of the half dozen patches that carry its information. With the block errors the four regions do not differ from each other.

The comparison between the rows of Table 5 is differential. It does not need the model to be right, only to be the same in every row.

6.4 What moves the index, and why we do not quote one

Table 7 lists what else moves the answer in the central configuration. The rows are changes of the index; the statistical error of the baseline is given for scale.

Angular scale. To ask at which scales two bands disagree we first marginalized the modes of one map at multipoles 2 to L and refitted, with the same cut in the K, Ka fit as a control fixed in advance, where nothing should move. The control failed: the K, Ka index moved toward flat in every region, by +0.08, +0.16, +0.11 and +0.37 at $L = 20$ (white CMB prior; E042). Read from the other side, with the pixel scale excluded and the index informed by multipoles up to 30, 20 or 12 only, the four-band index steepens (Table 7): by 0.017 to 0.055 at High Latitude, 0.017 to 0.032 in the Spur, 0.045 to 0.093 in the Galactic Center and 0.055 to 0.110 in the Plane. The estimator is sound: on simulations with known indices it is unbiased and calibrated (standardized errors of standard deviation 0.95 to 1.03 over 100 realizations at $N_{\rm side} = 8$), an independent review re-derived it and confirmed it on its own simulations, and the difference between the large-scale and the all-scale estimate of one simulated sky has zero mean and a spread of 0.9 to 1.1 times

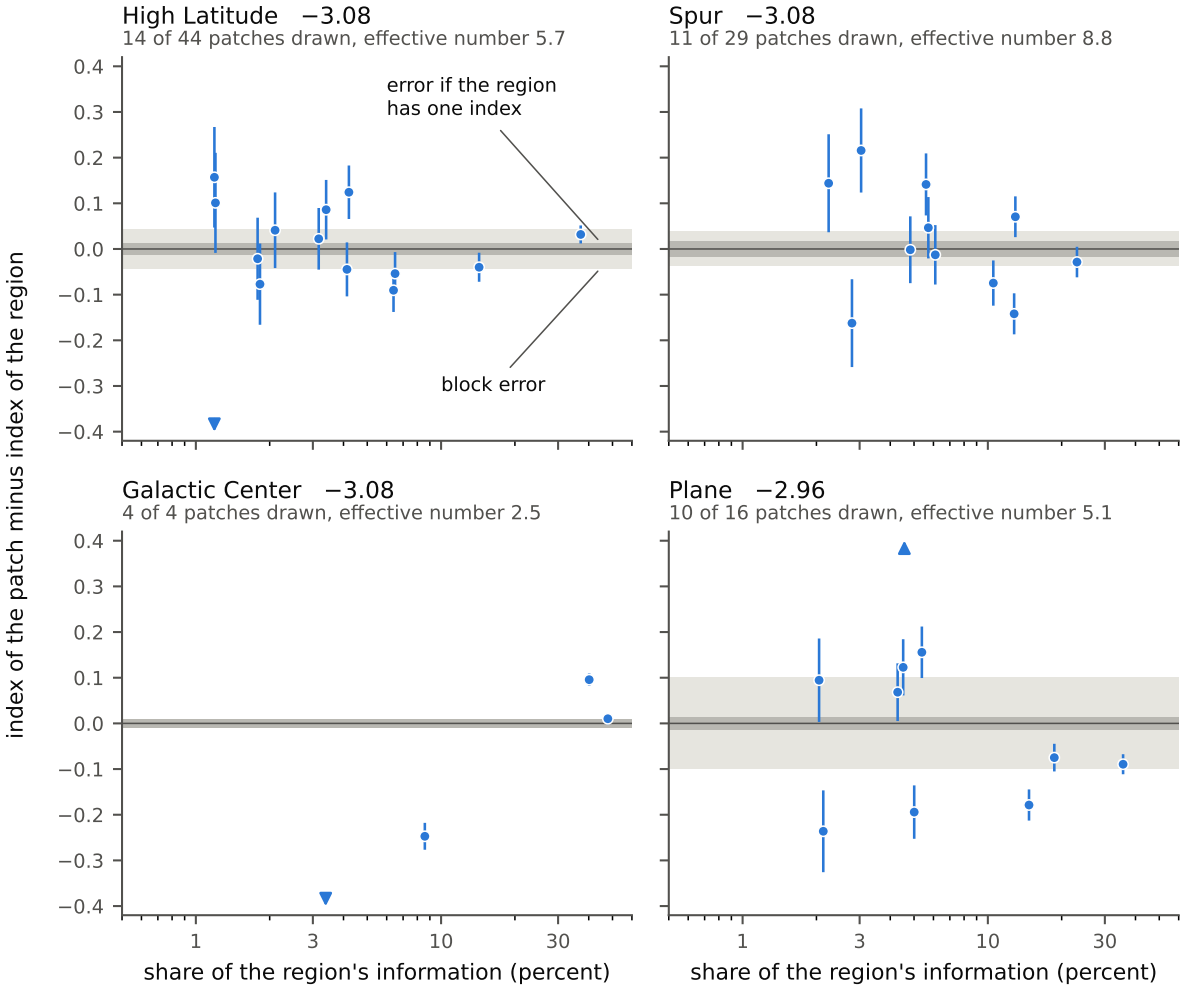


Figure 3: What the block errors of Table 5 are made of. For each patch of sky (a pixel of the $N_{\text{side}} = 2$ grid) that carries information on the index of a region: the index that the patch alone would move to in one Newton step from the fitted index of the region, with the error it would have under the model, against the share of the region's information that the patch carries. Patches with a one-step error above 0.12 are not drawn, and arrows mark values outside the frame. Dark band: the sandwich error of the region's index. Light band: its block error (none for the Galactic Center). The fitted index of the region is given beside its name. The one-step values are first-order and for display; the block error is computed from the scores, Eq. (24) (E040).

Table 7: Changes of the profile index of the WMAP K to V fit of Table 5 under analysis choices that have nothing to do with the estimator (E040, E042). "Index from $\ell \leq L$": K band and 353 GHz keep all their modes, which with the prior on the CMB fixes the amplitudes, and Ka, Q and V inform the index through their harmonic coefficients at multipoles 2 to L only. The window row compares two fits with the white CMB prior, and the last row is the K, Ka fit minus the four-band fit when both are read from multipoles up to 20. The flat-prior row repeats Table 5. The statistical error of the baseline is given for scale.

	High Latitude	Spur	Galactic Center	Plane
baseline: all scales, ΛCDM CMB prior	−3.082	−3.080	−3.081	−2.964
statistical error (sandwich)	0.012	0.016	0.009	0.013
flat amplitude prior at 30 GHz	+0.158	+0.090	+0.027	+0.045
no window matching (white CMB prior)	+0.007	+0.002	+0.009	+0.004
CMB prior: white, 0.2 μK	−0.012	−0.017	−0.006	−0.031
CMB prior: flat	+0.055	+0.025	+0.073	−0.024
index from $\ell \leq 30$	−0.017	−0.017	−0.045	−0.055
index from $\ell \leq 20$	−0.024	−0.031	−0.055	−0.084
index from $\ell \leq 12$	−0.055	−0.032	−0.093	−0.110
K and Ka only	+0.024	+0.012	+0.038	−0.002
dust index 1.48 in place of 1.53	−0.012	−0.008	−0.004	−0.027
dust index 1.58	+0.011	+0.007	+0.004	+0.025
no bandpass drift correction	−0.001	−0.001	−0.001	−0.001
353 GHz: all bolometers	+0.002	+0.001	+0.001	+0.003
delta bandpasses, nominal frequencies	−0.061	−0.057	−0.051	−0.076
K and Ka only, both fits from $\ell \leq 20$	+0.014	+0.024	+0.029	+0.009

$(\sigma^2_{\rm low} - \sigma^2_{\rm all})^{1/2}$ (E039). On the maps the differences at $L = 20$ are 2.8, 3.3, 8.2 and 10.4 times that number. The shifts are the same within 0.025 with the white CMB prior. So the dependence on scale is a property of these maps under this model, in every region, and not of the estimator or of the CMB prior. An index read from the three cuts is not stable either: the largest difference between two cuts is 2.8, 2.0, 5.4 and 6.7 times the standard deviation expected for nested fits.

We do not know the cause. One candidate is instrumental: the maps are pixel averages of skies that are not band-limited, structure below the pixel scale folds into it through each band's own beam, K band has the widest, and the window operator works to multipole 47 with Gaussian beams. Against it, the one simulation we have of sub-pixel structure leaves a residual of 0.002, though that sky had no compact sources. The other candidate is the sky: compact structure in the plane may be flatter than the diffuse emission, and an index that varies from patch to patch (previous subsection) makes an index from large scales and an index from all scales weight the sky differently. Fuskeland et al. (2014), with maps smoothed to a common 1° beam before any pixel averaging, find -2.98 ± 0.01 in the plane, close to our all-scale K, Ka value there (−2.97) and not to the large-scale one (−3.04 at $L = 20$); that favours the sky. The rows of Table 7 are sizes of an effect and not corrections.

Band subset. K and Ka alone give indices that differ from the four bands together by +0.024, +0.012, +0.038 and −0.002. The two fits share K and Ka, so the standard deviation of their difference under the model is $(\sigma^2_{\rm sub} - \sigma^2_{\rm full})^{1/2}$, and the differences are +2.5, +0.9, +5.2 and −0.2 of it; with both fits read from multipoles up to 20 they are +1.3, +1.8, +3.4 and +0.8. With the white CMB prior the shifts were larger and all of one sign (+0.032, +0.023, +0.041 and +0.015). What remains with the ΛCDM prior is a flatter K to Ka index at High Latitude and in the Galactic Center, where the model does not fit in any case. A spectrum that steepens with frequency would do this, and so would a dust model that predicts too much polarized dust at 41 and 61 GHz from 353 GHz.

Dust. The dust spectrum is fixed. Changing its index by 0.05 either way moves the synchrotron index by 0.012, 0.008, 0.004 and 0.026, linearly ($\partial\beta_\mathrm{s}/\partial\beta_\mathrm{d}$ = 0.24, 0.15, 0.08 and 0.52): one statistical standard deviation at High Latitude and two in the Plane, more than we had written down in advance for three of the four regions. With the uncertainty of the Planck value, 0.02, it is 0.005 at High Latitude and 0.010 in the Plane.

Bandpasses. Leaving out the correction for the drift of the WMAP band centres changes the index by 0.001, and taking the 353 GHz band of all bolometers in place of the polarization-sensitive ones by 0.001 to 0.003. Replacing the measured bandpasses by delta functions at the nominal WMAP frequencies steepens it by 0.051 to 0.076, 3.5 to 5.8 statistical standard deviations: the effective frequency of K band for a spectrum this steep is 22.4 GHz, not 22.8.

Poorly measured modes. The released K-band covariance has 20 eigenmodes with more than ten times the median variance. Projecting all of them out moves the K, Ka index by 0.002 to 0.015 (white CMB prior). This tests the modes the matrix knows about, not the claim that it understates them.

Taken together: for these maps and regions the profile index of the simple model is −3.08 at High Latitude, in the Spur and in the Galactic Center and −2.96 in the Plane. Its error is 0.01 to 0.02 if each region has one index and 0.04 to 0.10 from the scatter between patches. It moves by 0.01 to 0.07 with the prior on the CMB, by 0.02 to 0.11 with the angular scales that inform it, by up to 0.04 with the band subset, by 0.004 to 0.026 for 0.05 in the dust index, and by 0.05 to 0.08 if the bands are taken at their nominal frequencies. A number with an error bar that means what it says needs a model in which a region is more than one index and the scale dependence understood. That is a separate piece of work, and we do not quote an index here.

6.5 Planck 30 GHz

The WMAP K band and the Planck 30 GHz band are 0.22 apart in $\ln\nu$, and their polarization maps have been known to disagree on large scales since the 2015 Planck release. Weiland et al. (2022) found a K to 30 GHz index of −3.27 in the Fan region against −2.90 from K to Ka with the 2018 map and set 30 GHz aside. Watts et al. (2023) attribute the WMAP side to poorly measured modes of the K-band map and report that after a joint reprocessing of both data sets the two maps agree to within the noise at high latitude. We have only the official maps, and for them the question can be put in likelihood terms without a new claim about its cause.

For the PR4 30 GHz map the white-noise covariance of the map file is not enough at this resolution: the half-ring difference, averaged to $N_\mathrm{side} = 16$ with the same weights, has a χ^2 per element of 1.77 against it, with an excess of 1.7 at the pixel scale rising to between 3 and 7 at multipoles 2 to 6 (E038). We model the noise as $\mathsf{W}^{1/2}\mathsf{R}\mathsf{W}^{1/2}$, with W the white covariance and $\mathsf{R} = r\,1 + \mathsf{Y}\,\mathrm{diag}(A\,(\ell/10)^{-\gamma})\,\mathsf{Y}^\mathsf{T}$, $(r, A, \gamma) = (1.46, 1.16, 1.27)$ fitted to the whitened half-ring difference, which then passes 22 sub-tests with none beyond 3σ. The model is isotropic, fitted to one realization, and blind to anything common to the two halves. (The low-resolution covariance matrices described for PR4 by Planck Collaboration Int. LVII 2020 were not retrievable from the public mirrors when this was done.)

With this model, K band and PR4 30 GHz with 353 GHz do not fit one power law per region: the reduced χ^2 is 2.6, against 1.1 for K with Ka, and the fitted index is -3.43 ± 0.02 at high latitude where K with Ka gives -3.06 ± 0.02. With the 2018 map in place of PR4 and a noise model built the same way the reduced χ^2 is 6.0 and the fitted index is on the other side, −2.92. A fit this bad does not measure an index, and we do not interpret these two numbers beyond their disagreement.

A direct view needs no noise model for 30 GHz. We fit K, Ka, Q, V and 353 GHz, predict the 30 GHz map from the fitted amplitudes and indices, and subtract (Fig. 4); the same exercise with Ka predicted from K, Q, V and 353 GHz is the control. The residual of the 30 GHz map

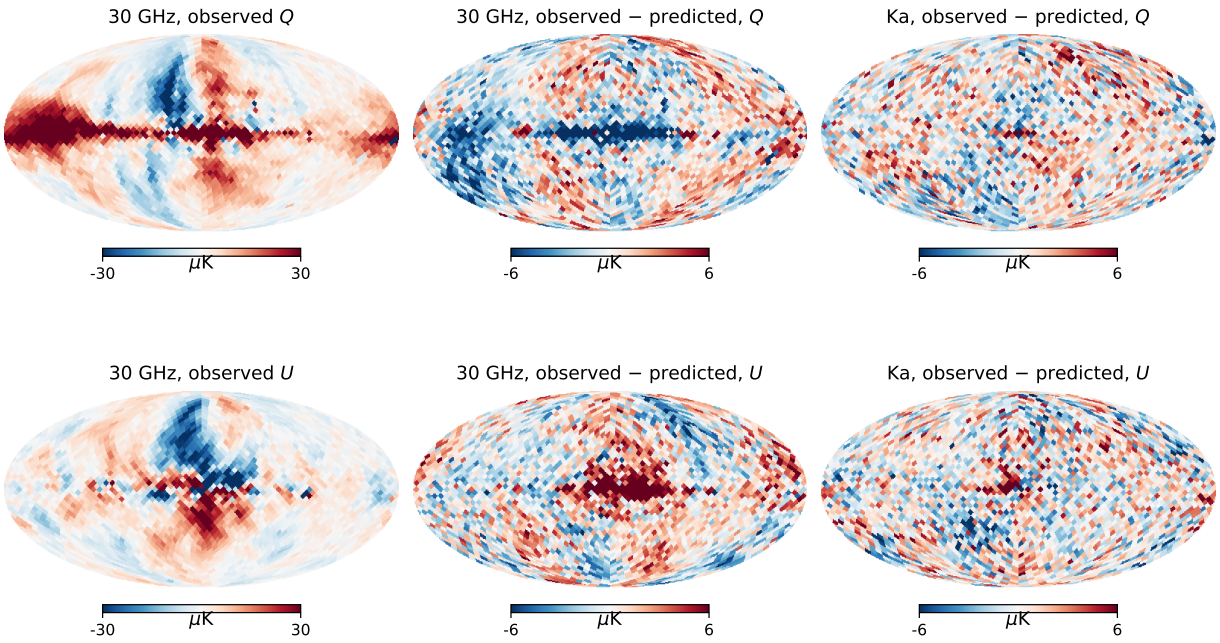


Figure 4: Left: the PR4 30 GHz Stokes Q and U maps at $N_{\rm side} = 16$. Centre: the same maps minus their prediction from a fit to WMAP K, Ka, Q, V and 353 GHz that uses no LFI band. Right: the same difference for WMAP Ka predicted from K, Q, V and 353 GHz, on the same colour scale (E042).

has 16 to 43 times the power of its noise model at multipoles 2 to 6, in B as much as in E; the control has 2 to 4 times its own, part of which is the noise of the prediction. A residual with as much B as E that follows the inner Galactic plane and has large-scale structure at intermediate latitudes is map-level structure, not sky signal assigned to the wrong index. Removing the 20 poorly measured K-band modes changes the K, 30 index by 0.03. We conclude only this: at a statistical precision of 0.01 to 0.02 in the index, the official 30 GHz polarization maps cannot be combined with the official WMAP K to V maps in one power-law fit, and an index from a band set that mixes them is set by that inconsistency. Further exploratory comparisons of the PR4 and 2018 maps are in the record (E042).

6.6 A band set without K band at a 30 GHz reference

Sect. 4 says the volume term is largest when the reference frequency is far from the pivot, the amplitudes are many and other free amplitudes move the pivot. One configuration in use has all three: the band set 30, 44, 70 GHz, Ka, Q, V and 353 GHz without K band, a synchrotron amplitude defined at 30 GHz, and a CMB amplitude without a prior (Svalheim et al., 2023). On the BeyondPlanck maps themselves an earlier experiment of ours found the flat-prior marginal of that band set shifted from the profile optimum by +0.27 in the Spur and +0.13 in the Plane at $N_{\rm side} = 16$, and at the edge of its grid at high latitude (E022). Table 8 repeats the calculation on the PR4 and WMAP maps.

With flat priors on every amplitude and a 30 GHz reference, the optimum of the index marginal for the High Latitude region is at the upper edge of the allowed range, more than a unit of index flatter than the profile optimum of the same likelihood; the Spur sits at -2.31 and the Plane at -2.48. Defined at 23 GHz the same marginal puts the High Latitude region at -4.04. Nothing about the data changes between the three rows. [With the white CMB prior: B2.]

Two published analyses in this configuration constrain the high-latitude index with a prior $\mathcal{N}(-3.3, 0.1^2)$ instead of fitting it: Svalheim et al. (2023) because of "very strong degeneracies with respect to instrumental systematic effects" there, and Watts et al. (2023), who report that

Table 8: The band set 30, 44, 70 GHz (PR4), Ka, Q, V (WMAP) and 353 GHz, without K band, with flat priors on the synchrotron, dust and CMB amplitudes, at $N_{\rm side} = 16$ in the four regions (E040). The profile optimum, and the optimum of the flat-prior marginal with the synchrotron amplitude defined at 23 and at 30 GHz. The allowed range of the index is $[-4.6, -1.6]$.

	High Latitude	Spur	Galactic Center	Plane
profile likelihood	-2.70 ± 0.04	-2.92 ± 0.05	-2.82 ± 0.03	-2.78 ± 0.04
flat priors, reference 23 GHz	-4.04	-3.43	-2.89	-3.00
flat priors, reference 30 GHz	-1.60 (bound)	-2.31	-2.75	-2.48

when the polarized synchrotron amplitude and the calibration are both fitted freely the high-latitude index tends to "unreasonably flat values", $\beta_{\rm s} \lesssim -2.5$, which they attribute to unmodelled systematics. The volume term of this configuration has that sign and is large. We do not claim more. The published high-latitude values are prior draws, so there is no fitted number to compare with; the calibration is outside a maps-only calculation; and the same calculation predicts pulls of $+0.6$ and $+0.3$ in the Spur and the Plane on our maps, where the published values, fitted with the prior, are -3.17 ± 0.06 and -3.03 ± 0.07, so the volume term alone does not describe those chains. Their maps, pixelization, amplitude priors and mixing all differ from this calculation. The point for this paper is the size: in a configuration that is in use, the flat-prior convention moves a region index by tenths to more than a unit, and the profile likelihood on the same maps moves it by nothing. (The profile optimum of this band set is itself unstable, -2.70 with a flat CMB prior and -2.99 with the white one at high latitude, which is the inconsistency of Sect. 6.5 seen from another side.)

7 Limits: when a proper amplitude prior is needed

The profile likelihood settles the question for a handful of indices that the data determine. It does not settle it for an index field with a hierarchical prior at low signal to noise, where a proper amplitude prior is needed, and there the prior becomes a model component with failure modes of its own. We list what we established, as limits on what the earlier sections recommend. All of it is on simulations.

A scale shared across index regions on a scale map. If the zero-mean prior takes its scale from a map, $a_p \sim \mathcal{N}(0, (s\, g_p)^2)$ with one s for the sky, a scale map that is wrong differently in different regions biases their indices with a sign that again depends on the reference frequency. On the sky of Table 2, a map that is 30 percent high at $|b| > 30^\circ$ moves the high-latitude indices by -0.026 at a 23 GHz reference, $+0.029$ at 30 GHz and 0.000 at the pivot (E036). The effect is coherent over a zone and of the order of its pooled error. One free scale per index region does not have it.

Template-tied priors in intensity. For positive intensity amplitudes we implemented $a_p \sim \mathcal{N}(\mu\, g_p, (s\, g_p)^2)$ with a template g and inferred μ and s per component. This prior passes simulation-based calibration on data drawn from itself (600 replicates, E016) and gives calibrated index fields down to index-pixel signal-to-noise ratios below one. It is exact only for its own template-error law. A template error of the same size with a different conditional law, such as a template that is the true amplitude divided by a 30 percent scatter, shifts the faint-sky dust index by 0.11 in the exact target and by 0.13, four posterior standard deviations, in the chain (E021): regressing amplitudes on a noisy template is itself an errors-in-variables problem. At high signal to noise with a template built from the data and a sky outside the model class (PySM d1s1), the posterior runs away: the index removes the component from all bands but one because that lowers the prior cost by more than it raises χ^2 (E020). At the white-noise depth of the WMAP

and LFI intensity maps with a 30 percent template error the field mean is within 0.02 of the truth (-3.017 against -3.000, five posterior standard deviations), but the inferred field does not beat a constant index in pixel error on that sky, and the field hyperparameters fail a four-chain split-$\hat{R}$ criterion (1.11 to 1.42) although every chain passes its own convergence check (E032).

What follows. An index field with inferred smoothness is a different paper. The conclusion for this one is narrower: where the question is one index per region and the region is measured, use the profile likelihood with the debiased sandwich; do not use a flat amplitude prior unless the reference frequency is the pivot; and do not share one amplitude scale across index regions on a template. None of this makes the index right if the model of the sky or of the instrument is wrong, which is the subject of the last part of Sect. 6.

8 Implementation and validation

Commander4. In our branch of Commander4 the spectral-index Metropolis step of the per-pixel path takes a `likelihood` option. `marginal` evaluates Eq. (4) and `ridge` Eq. (5) in the acceptance, with the amplitudes integrated or maximized analytically per pixel in a compiled kernel; `conditional` is the Gibbs step at fixed amplitudes; `resolve_chisq` keeps the move of Sect. 3 with a warning; `auto` resolves to `ridge` where the kernel applies and to `conditional` otherwise. Each exact option is tested against a grid evaluation of its target in unit tests. End to end, a seven-rank run on simulated maps matched the exact profile posterior within 0.03 posterior standard deviations with width ratios 0.94 and 0.97 (E006); in that realization the target itself sat 1.9 and 3.7 standard deviations from the input indices, which a 40-seed repetition showed to be scatter of the harness at its noise level and not an offset (E037). On real nside-64 intensity maps the `ridge` chain matched an independent grid evaluation of its target with shifts of -0.06 ± 0.04 and $+0.06 \pm 0.05$ standard deviations and width ratios 0.96 and 0.94 for the two indices (E014). On the same maps the conditional Gibbs step has a standard deviation 180 times smaller than the marginal one for the synchrotron index and did not move in 2000 iterations, while the collapsed chains mix in tens of iterations: collapsing the amplitudes is needed for mixing as well as for the target. Simulation-based calibration of the ridge and marginal modes inside Commander4 itself has not been run.

The low-resolution module. The analysis of Sect. 6 uses a separate, small dense implementation (`nextgen/lowres`): band-averaged spectra from the instrument bandpasses, the collapsed objectives of Sect. 2 for several coupled regions with dense noise covariances, analytic gradients, and J and T of Eq. (6). At $N_{\mathrm{side}} = 16$ with three components the amplitude system has 18 432 unknowns, and a profile fit of four regions takes three to five minutes on a desktop machine. Unit tests check the gradients against finite differences, the diagonal-noise closed forms, J and T against $\sum uV$ and $\sum V$, the sandwich against the scatter of 300 fits, the plug-in bias and its correction, and the identity of Eq. (14) to nine digits. The bandpass readers reproduce the published LFI colour-correction ratios to 3×10^{-4} and the HFI unit conversions to 0.3 percent (E038), and the per-radiometer effective frequencies of the WMAP team's own routine to 0.005 GHz (E039). An independent statistical review of the module, by a reviewer who derived the targets from the stated model before reading our derivations, found no code defect; it corrected the response formula of Sect. 4 (the curvature of the marginal in place of the profile information), the condition under which the profile score is unbiased when a component carries a prior (App. A), and a docstring, and it measured the limits of the plug-in error quoted in Sect. 5 (E039, independent review).

A numerical hazard. During this work three processes crashed with corrupted heaps. The cause was not in our code. On the machine used (Apple M4, macOS 27.0, a pre-release build)

the Cholesky routine of the Accelerate framework, called with the upper-triangle option, writes between 0.25 and 0.77 MB past an internal buffer for matrix orders of about 1500 to 6000, and SciPy maps `cho_factor(lower=True)` on a row-ordered array to that option. The factor it returns is correct, so no test of values can see the fault; what lies behind the buffer in memory is overwritten. We found it with the system's guard allocator and measured it with a malloc canary that is now part of the repository, switched the module to the other triangle (same cost, same numbers), and added a test that runs the likelihood under the canary (`E043`). The fits at $N_{\mathrm{side}} = 16$ (orders 12 288 and 18 432) were not exposed: none of several hundred factorizations of the baseline fit overran, and its output is reproduced in every printed digit. The simulations at $N_{\mathrm{side}} = 8$ (orders 1536 and 4608) were exposed, and every one of them was run again on the sound path under the canary and compared with the original output as printed. For the 200 realizations of Table 4 all 1480 printed numbers agree to the last digit, as do the 40 realizations with errors from each data set, and no run made since has touched the canary. We report this because the failure is silent, because dense low-resolution likelihoods of exactly this size are common in this field, and because a result that passes every statistical test can still have been computed next to an overwritten buffer.

9 Discussion and conclusions

What to do. For a spectral index shared by many pixels with linear amplitudes and known noise: (i) evaluate the profile likelihood or a proper marginal directly in the acceptance step, as the marginal and ridge samplers of Andersen et al. (2023) do, never the χ^2 at freshly drawn amplitudes; (ii) do not use a flat amplitude prior unless the amplitudes are defined at the pivot frequency of the band set, which depends on the other components and their priors, and say which reference frequency was used; (iii) quote the profile optimum with the debiased sandwich error, Eq. (22), and check $\mathsf{T}/\mathsf{J}_{\mathrm{c}}$; (iv) if a prior is put on a component that does not depend on the index, state the sensitivity of the index to that prior.

What the real maps show. On the WMAP K to V polarization maps, changing nothing but the frequency at which a flat amplitude prior is stated, from 23 to 30 GHz, moves the high-latitude index by 0.17 at $N_{\mathrm{side}} = 16$, 12.6 times its statistical error; with a flat prior on the CMB as well the move is 0.8; for a band set without K band it is more than a unit. In every case that we could check (four regions, three CMB priors, shifts from 0.03 to 0.8) the first-order identity predicts the move to 10 percent. The profile likelihood and the zero-mean prior do not depend on the convention. That comparison is differential and does not rest on the sky model being right. The prior on the CMB matters in a second way: a white prior with the right total variance and the wrong shape leaves the index unbiased on ΛCDM skies and understates its error by a quarter to a third.

What they do not show. We have not measured the polarized synchrotron index. The profile index of the simple model is near -3.08 at high latitude, in the North Polar Spur and in the Galactic Center and -2.96 in the Galactic plane, with errors of 0.01 to 0.02 if each region has one index. At high latitude and in the plane it does not: the scatter of the score over patches of sky gives errors of 0.04 and 0.10 there (and 0.04 in the Spur, where it is within what the model allows), half a dozen patches carry the information of each region, and with those errors the regions do not differ. The index also moves by 0.02 to 0.11 with the angular scales that inform it, in every region and by several times what the model allows for nested fits; by 0.01 to 0.07 with the prior on the CMB; by up to 0.03 for a change of 0.05 in the dust index, which is fixed; and by up to 0.04 with the band subset. The model misfits in the plane and fails in the Galactic Center; and the official 30 GHz maps do not fit one power law with K band. The spread of published values for this sky, -2.9 to -3.3, is of the size of these effects and of the convention

effect together, not of the statistical errors. A measurement needs a model in which a region is more than one index, the scale dependence understood, and simulations of each added step.

What we did not do. All results are for maps, at low resolution, with one index per region and Gaussian noise of known covariance. We did not treat time-ordered data, gain or bandpass sampling, or an index field with inferred smoothness, for which a proper amplitude prior is needed and has the failure modes of Sect. 7. The account of the real-map chain in Sect. 3 uses a factor measured from the log of the same chain, and the forward test has not been run. Simulation-based calibration of the exact modes inside Commander4 itself has not been run; they are tested against grids of their targets. The estimator that reads the index from large scales only is unbiased and calibrated on simulations in which the model is true; that says nothing about what it does with whatever makes the real maps depend on scale. The noise model for 30 GHz is empirical and isotropic. The block errors treat patches of sky as exchangeable and independent, rest on five to nine effective patches per region, are too small if the index varies coherently over several patches, and have to be read against simulations with the covariances of the fit. The ΛCDM prior describes multipoles up to 47; the power above that changes the indices by less than 0.001.

Reproducibility. Every number in this paper comes from a script in the repository and an experiment record that states the question, the predictions written before the run, the command and the outcome, including the predictions that failed. The records are cited inline.

A Score, curvature and sandwich for the profile likelihood

Let $\Pi = \mathsf{M}\mathsf{A}^{-1}\mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}$ be the projector onto the model space, so that $\mathsf{P} = \mathsf{N}^{-1}(1-\Pi)$, $\mathsf{P}\mathsf{M} = 0$ and $\chi^2_{\min} = d^{\mathsf{T}}\mathsf{P}d$. With $\mathsf{M}_r = \partial\mathsf{M}/\partial\beta_r$ and $\mathsf{P}_r = \partial\mathsf{P}/\partial\beta_r$, differentiating $\mathsf{P}\mathsf{M} = 0$ gives $\mathsf{P}_r\mathsf{M} = -\mathsf{P}\mathsf{M}_r$. The score is $s_r = -\frac{1}{2}d^{\mathsf{T}}\mathsf{P}_r d$. For $d = \mathsf{M}a + n$:

$$s_r = -\tfrac{1}{2}a^{\mathsf{T}}\mathsf{M}^{\mathsf{T}}\mathsf{P}_r\mathsf{M}a - a^{\mathsf{T}}\mathsf{M}^{\mathsf{T}}\mathsf{P}_r n - \tfrac{1}{2}n^{\mathsf{T}}\mathsf{P}_r n = (\mathsf{M}_r a)^{\mathsf{T}}\mathsf{P}\, n - \tfrac{1}{2}n^{\mathsf{T}}\mathsf{P}_r n, \tag{25}$$

because $\mathsf{M}^{\mathsf{T}}\mathsf{P}_r\mathsf{M} = -\mathsf{M}^{\mathsf{T}}\mathsf{P}\mathsf{M}_r = 0$. The first term has zero mean. The second has mean $-\frac{1}{2}\mathrm{tr}(\mathsf{P}_r\mathsf{N}) = \frac{1}{2}\,\partial\,\mathrm{tr}(\Pi)/\partial\beta_r = 0$, since $\mathrm{tr}\Pi = k$ for every β. So $\mathrm{E}[s_r] = 0$ for any true amplitudes: the profile score is unbiased.

The variance of the first term is $(\mathsf{M}_r a)^{\mathsf{T}}\mathsf{P}\mathsf{N}\mathsf{P}(\mathsf{M}_{r'}a) = (\mathsf{M}_r a)^{\mathsf{T}}\mathsf{P}(\mathsf{M}_{r'}a) = \mathsf{J}_{rr'}$, using $\mathsf{P}\mathsf{N}\mathsf{P} = \mathsf{P}$. The second term is a quadratic form in Gaussian noise, uncorrelated with the first, with covariance $\frac{1}{2}\mathrm{tr}(\mathsf{P}_r\mathsf{N}\mathsf{P}_{r'}\mathsf{N}) = \frac{1}{2}\mathrm{tr}(\Pi_r\Pi_{r'})$. Writing $\mathsf{G} = \mathsf{A}^{-1}\mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}$,

$$\Pi_r = (1-\Pi)\mathsf{M}_r\mathsf{G} + \mathsf{M}\mathsf{A}^{-1}\mathsf{M}_r^{\mathsf{T}}\mathsf{N}^{-1}(1-\Pi), \tag{26}$$

and since $\mathsf{G}(1-\Pi) = 0$ and $(1-\Pi)\mathsf{M} = 0$ the two squared terms vanish in the trace and the cross terms give

$$\tfrac{1}{2}\mathrm{tr}(\Pi_r\Pi_{r'}) = \mathrm{tr}\big(\mathsf{A}^{-1}\mathsf{M}_r^{\mathsf{T}}\mathsf{P}\,\mathsf{M}_{r'}\big) = \mathsf{T}_{rr'}. \tag{27}$$

The expected curvature follows from differentiating the score once more: $\mathrm{E}[\partial^2(\frac{1}{2}\chi^2_{\min})/\partial\beta_r\partial\beta_{r'}] = \mathsf{J}_{rr'}$. The estimator $\hat{\beta}$ that solves $s(\hat{\beta}) = 0$ therefore has covariance $\mathsf{J}^{-1}(\mathsf{J}+\mathsf{T})\mathsf{J}^{-1}$ to leading order, Eq. (19).

A Gaussian prior on a component whose spectrum does not depend on the index (the CMB with variance σ_{c}^2 per element) fits in by minimizing over that component first: the result is the flat-prior problem for the remaining components with noise covariance $\mathsf{N} + \mathsf{C}\,\Lambda^{-1}\mathsf{C}^{\mathsf{T}}$, with C the columns of the constrained component and Λ its prior precision. Everything above then holds in expectation over noise and over skies of that component drawn from the prior, with its sample variance counted in $\mathsf{J}+\mathsf{T}$. It does not hold for a prior variance that differs from the true one. If the true variance is κ^2 times the prior variance, the expected gradient of $\frac{1}{2}\chi^2_{\min}$ is $(\kappa^2-1)\,t$ (the expected score is minus that) with $t_r = \mathrm{tr}(\Lambda\mathsf{A}^{-1}\mathsf{M}_r^{\mathsf{T}}\mathsf{B}^{\mathsf{T}})$, $\mathsf{B} = \mathsf{A}^{-1}\mathsf{M}^{\mathsf{T}}\mathsf{N}^{-1}$, A containing the

prior precision, and the optimum is biased by $-(\kappa^2 - 1)\,\mathsf{J}^{-1}t$ to first order. A single factor κ^2 describes a prior that is wrong in size; a prior that is wrong in shape, as a white prior is for the CMB, needs the calculation mode by mode. Because t_r is a sum over elements of one sign, the bias in units of the standard deviation grows as the square root of the number of elements. We print this sensitivity with every fit that uses such a prior.

For the fitted amplitudes $\hat{a} = a+e$ with $\mathrm{Cov}(e) = \mathsf{A}^{-1}$, $\mathrm{E}[(\mathsf{M}_r\hat{a})^\mathsf{T}\mathsf{P}(\mathsf{M}_{r'}\hat{a})] = \mathsf{J}_{rr'}+\mathrm{tr}(\mathsf{A}^{-1}\mathsf{M}_r^\mathsf{T}\mathsf{P}\mathsf{M}_{r'}) = \mathsf{J}_{rr'} + \mathsf{T}_{rr'}$, which gives Eq. (22).

For one power-law component, delta bands and diagonal noise, $\mathsf{M}_r a$ has elements $a_p m_b \ln(\nu_b/\nu_{\mathrm{ref}})$, the projector removes the component of $\ln\nu$ along the amplitude direction, and $\mathsf{J} = \sum_p a_p^2 A_p V_p$, $\mathsf{T} = \sum_p V_p$ with $A_p = \sum_b w_{pb} m_b^2$ and V_p the weighted variance of $\ln\nu_b$ under the weights $w_{pb}m_b^2$, Eq. (7).

References

Alonso, D., Dunkley, J., Naess, S., & Thorne, B. 2017, Phys. Rev. D, 95, 043504 (arXiv:1608.00551)

Alquier, P., Friel, N., Everitt, R., & Boland, A. 2016, Stat. Comput., 26, 29 (arXiv:1403.5496)

Andersen, K. J., Herman, D., Aurlien, R., et al. 2023, A&A, 675, A13 (BeyondPlanck XIII; arXiv:2201.08188)

Andrieu, C., & Roberts, G. O. 2009, Ann. Statist., 37, 697 (arXiv:0903.5480)

Ball, R. C., Fink, T. M. A., & Bowler, N. E. 2003, Phys. Rev. Lett., 91, 030201 (arXiv:cond-mat/0301179)

Bardenet, R., Doucet, A., & Holmes, C. 2017, J. Mach. Learn. Res., 18, 1 (arXiv:1505.02827)

Beaumont, M. A. 2003, Genetics, 164, 1139

Bell, R. M., & McCaffrey, D. F. 2002, Survey Methodology, 28, 169

Bennett, C. L., Larson, D., Weiland, J. L., et al. 2013, ApJS, 208, 20 (arXiv:1212.5225)

Cameron, A. C., & Miller, D. L. 2015, J. Human Resources, 50, 317

Carter, A. V., Schnepel, K. T., & Steigerwald, D. G. 2017, Rev. Econ. Stat., 99, 698

Ceperley, D. M., & Dewing, M. 1999, J. Chem. Phys., 110, 9812

Dickinson, C., Eriksen, H. K., Banday, A. J., et al. 2009, ApJ, 705, 1607 (arXiv:0903.4311)

Dunkley, J., Spergel, D. N., Komatsu, E., et al. 2009, ApJ, 701, 1804 (arXiv:0811.4280)

Engel, B. 1985, Kwantitatieve Methoden, 19, 5

Eriksen, H. K., Jewell, J. B., Dickinson, C., et al. 2008, ApJ, 676, 10 (arXiv:0709.1058)

Fuskeland, U., Wehus, I. K., Eriksen, H. K., & Næss, S. K. 2014, ApJ, 790, 104 (arXiv:1404.5323)

Huber, P. J. 1967, in Proc. Fifth Berkeley Symp. Math. Statist. Probab., Vol. 1, 221–233

Jarosik, N., Barnes, C., Greason, M. R., et al. 2007, ApJS, 170, 263 (arXiv:astro-ph/0603452)

Jarosik, N., Bennett, C. L., Dunkley, J., et al. 2011, ApJS, 192, 14 (arXiv:1001.4744)

Jew, L., & Grumitt, R. D. P. 2020, MNRAS, 495, 578 (arXiv:1907.11426)

Lewis, A., Challinor, A., & Lasenby, A. 2000, ApJ, 538, 473 (arXiv:astro-ph/9911177)

Liang, K.-Y., & Zeger, S. L. 1986, Biometrika, 73, 13

Medina-Aguayo, F. J., Lee, A., & Roberts, G. O. 2016, Stat. Comput., 26, 1187 (arXiv:1503.07066)

Neyman, J., & Scott, E. L. 1948, Econometrica, 16, 1

Nicholls, G. K., Fox, C., & Watt, A. M. 2012, arXiv:1205.6857

O'Neill, P. D., Balding, D. J., Becker, N. G., Eerola, M., & Mollison, D. 2000, Appl. Statist., 49, 517

Planck Collaboration Int. LVII. 2020, A&A, 643, A42 (arXiv:2007.04997)

Planck Collaboration VI. 2020, A&A, 641, A6 (arXiv:1807.06209)

Planck Collaboration X. 2016, A&A, 594, A10 (arXiv:1502.01588)

Planck Collaboration XI. 2020, A&A, 641, A11 (arXiv:1801.04945)

Stompor, R., Leach, S., Stivoli, F., & Baccigalupi, C. 2009, MNRAS, 392, 216 (arXiv:0804.2645)

Svalheim, T. L., Andersen, K. J., Aurlien, R., et al. 2023, A&A, 675, A14 (BeyondPlanck XIV; arXiv:2011.08503v2)

Watts, D. J., Basyrov, A., Eskilt, J. R., et al. 2023, A&A, 679, A143 (arXiv:2303.08095)

Weiland, J. L., Addison, G. E., Bennett, C. L., Halpern, M., & Hinshaw, G. 2022, ApJ, 936, 24 (arXiv:2203.11445)

White, H. 1982, Econometrica, 50, 1–25